

OpenShift

User Guide

Using OpenShift to manage your applications in the cloud



OPENSIFT

OpenShift User Guide

Using OpenShift to manage your applications in the cloud

Edition 2.2.15

Copyright © 2012 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at <http://creativecommons.org/licenses/by-sa/3.0/>. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, JBoss, MetaMatrix, Fedora, the Infinity Logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux® is the registered trademark of Linus Torvalds in the United States and other countries.

Java® is a registered trademark of Oracle and/or its affiliates.

XFS® is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

All other trademarks are the property of their respective owners.

1801 Varsity Drive
Raleigh, NC 27606-2072 USA
Phone: +1 919 754 3700
Phone: 888 733 4281
Fax: +1 919 754 3701

This guide provides an introduction to OpenShift and documents its application management functions.

Preface	v
1. Document Conventions	v
1.1. Typographic Conventions	v
1.2. Pull-quote Conventions	vii
1.3. Notes and Warnings	vii
2. Getting Help	viii
2.1. Do You Need Help?	viii
2.2. We Need Feedback!	viii
 1. OpenShift Architecture Overview	 1
1.1. Platform Overview	1
1.2. System Resources and Application Containers	1
1.3. OpenShift Applications	2
1.3.1. Using Scaled Applications	4
1.4. Cartridges	5
 2. OpenShift Web Interface	 7
2.1. Accessing the OpenShift Management Console	7
2.2. Managing Your OpenShift Account	7
2.2.1. Changing Your Password	7
2.2.2. Resetting Your Password	8
2.2.3. Changing Namespaces	8
2.2.4. Managing SSH Keys	9
2.3. Creating Applications	10
2.3.1. Creating Basic Applications	10
2.3.2. Creating Preconfigured Applications	11
2.3.3. Cloning Application Files	12
2.4. Monitoring Application Resources	12
2.5. Editing and Deploying Applications	15
2.5.1. Preparing Your Application for Deployment	15
2.5.2. Deploying Your Application to the Cloud	16
2.6. Deleting Applications	16
2.6.1. Deleting Remote Application Data	16
2.6.2. Deleting Local Application Data	17
2.7. Using Cartridges	17
2.7.1. Adding Cartridges	17
 3. OpenShift Command Line Interface	 19
3.1. Secure Shell Authentication	19
3.1.1. Resolving Authentication Issues	19
3.1.2. Managing Multiple SSH Keys	20
3.2. Working With Domains	21
3.2.1. Creating a Domain	21
3.2.2. Altering a Domain	22
3.2.3. Deleting a Domain	23
3.3. Viewing User Information	24
3.4. Creating Applications	24
3.4.1. Introduction	24
3.4.2. Creating Applications	25
3.5. Editing and Deploying Applications	27
3.5.1. Preparing Your Application for Deployment	27
3.5.2. Deploying Your Application to the Cloud	28
3.5.3. Hot Deploying Applications	28
3.5.4. Deploying JBoss Applications	29
3.6. Using Cartridges with the CLI	32

3.6.1. Adding Cartridges with the CLI	32
3.6.2. Managing Cartridges	33
3.6.3. Working With Database Cartridges	33
3.7. Creating Application Snapshots	37
3.8. Managing Applications	38
3.8.1. Using Application Management Commands	38
3.8.2. Managing Applications in a Shell Environment	38
3.8.3. Managing Applications in a Secure Shell Environment on Windows	39
3.8.4. Using the Jenkins Embedded Build System	50
3.8.5. Managing Application Builds with Jenkins	51
3.8.6. Using Environment Variables	55
3.8.7. Using Node.js	56
3.8.8. Scheduling Timed Jobs with Cron	58
3.8.9. Sending and Receiving E-mail	60
4. Application Maintenance, Monitoring and Troubleshooting	61
4.1. Monitoring Applications with the MongoDB Monitoring Service (MMS)	61
4.1.1. Setting up MMS	61
4.1.2. Monitoring Your Applications with MMS	62
4.2. Managing Your Application Disk Space	64
4.2.1. The Disk Space Cleanup Tool	64
4.3. Troubleshooting JBoss Applications	65
4.3.1. Using Thread Dumps to Troubleshoot JBoss Applications	65
4.3.2. Inspecting Server, Boot and Other Log Files	66
4.4. Performing Application Maintenance from Your Workstation	67
4.4.1. Port Forwarding	67
5. Storage Management	69
5.1. Backing up and Restoring Configuration and User Data	69
5.1.1. Creating Snapshots	69
5.1.2. Restoring Snapshots	70
A. Revision History	71
Index	79

Preface

OpenShift is an enterprise-class *Platform-as-a-Service (PaaS)*. OpenShift provides enterprise developers with a wide selection of programming languages and frameworks including Java, Ruby, PHP, Perl, Python, and Node.js. It also provides integrated developer tools to support the application lifecycle, including Eclipse integration, JBoss Developer Studio, Jenkins, Maven, and GIT. OpenShift uses an open source ecosystem to provide key platform services for mobile applications (Appcelerator), NoSQL services (MongoDB), SQL services (Postgres, MySQL), and more. JBoss provides an enterprise-class middleware platform for Java applications, providing support for Java EE6 and integrated services such as transactions and messaging, which are critical for enterprise applications.

The foundation of the OpenShift platform is Red Hat Enterprise Linux, which provides a secure and scalable multi-tenant operating system to address the needs of enterprise-class applications as well as providing integrated application runtimes and libraries.

This guide provides an overview of the OpenShift architecture, and continues with detailed descriptions and procedures for how to use the OpenShift client tools to create namespaces, to create, deploy, and update applications, to monitor and troubleshoot applications, and to perform remote application maintenance and other tasks.

1. Document Conventions

This manual uses several conventions to highlight certain words and phrases and draw attention to specific pieces of information.

In PDF and paper editions, this manual uses typefaces drawn from the *Liberation Fonts*¹ set. The Liberation Fonts set is also used in HTML editions if the set is installed on your system. If not, alternative but equivalent typefaces are displayed. Note: Red Hat Enterprise Linux 5 and later includes the Liberation Fonts set by default.

1.1. Typographic Conventions

Four typographic conventions are used to call attention to specific words and phrases. These conventions, and the circumstances they apply to, are as follows.

Mono-spaced Bold

Used to highlight system input, including shell commands, file names and paths. Also used to highlight keycaps and key combinations. For example:

To see the contents of the file **my_next_bestselling_novel** in your current working directory, enter the **cat my_next_bestselling_novel** command at the shell prompt and press **Enter** to execute the command.

The above includes a file name, a shell command and a keycap, all presented in mono-spaced bold and all distinguishable thanks to context.

Key combinations can be distinguished from keycaps by the hyphen connecting each part of a key combination. For example:

Press **Enter** to execute the command.

¹ <https://fedorahosted.org/liberation-fonts/>

Press **Ctrl+Alt+F2** to switch to the first virtual terminal. Press **Ctrl+Alt+F1** to return to your X-Windows session.

The first paragraph highlights the particular keycap to press. The second highlights two key combinations (each a set of three keycaps with each set pressed simultaneously).

If source code is discussed, class names, methods, functions, variable names and returned values mentioned within a paragraph will be presented as above, in **mono-spaced bold**. For example:

File-related classes include **filesystem** for file systems, **file** for files, and **dir** for directories. Each class has its own associated set of permissions.

Proportional Bold

This denotes words or phrases encountered on a system, including application names; dialog box text; labeled buttons; check-box and radio button labels; menu titles and sub-menu titles. For example:

Choose **System** → **Preferences** → **Mouse** from the main menu bar to launch **Mouse Preferences**. In the **Buttons** tab, click the **Left-handed mouse** check box and click **Close** to switch the primary mouse button from the left to the right (making the mouse suitable for use in the left hand).

To insert a special character into a **gedit** file, choose **Applications** → **Accessories** → **Character Map** from the main menu bar. Next, choose **Search** → **Find...** from the **Character Map** menu bar, type the name of the character in the **Search** field and click **Next**. The character you sought will be highlighted in the **Character Table**. Double-click this highlighted character to place it in the **Text to copy** field and then click the **Copy** button. Now switch back to your document and choose **Edit** → **Paste** from the **gedit** menu bar.

The above text includes application names; system-wide menu names and items; application-specific menu names; and buttons and text found within a GUI interface, all presented in proportional bold and all distinguishable by context.

Mono-spaced Bold Italic or ***Proportional Bold Italic***

Whether mono-spaced bold or proportional bold, the addition of italics indicates replaceable or variable text. Italics denotes text you do not input literally or displayed text that changes depending on circumstance. For example:

To connect to a remote machine using ssh, type **ssh *username@domain.name*** at a shell prompt. If the remote machine is **example.com** and your username on that machine is john, type **ssh john@example.com**.

The **mount -o remount *file-system*** command remounts the named file system. For example, to remount the **/home** file system, the command is **mount -o remount /home**.

To see the version of a currently installed package, use the **rpm -q *package*** command. It will return a result as follows: ***package-version-release***.

Note the words in bold italics above — *username*, *domain.name*, *file-system*, *package*, *version* and *release*. Each word is a placeholder, either for text you enter when issuing a command or for text displayed by the system.

Aside from standard usage for presenting the title of a work, italics denotes the first use of a new and important term. For example:

Publican is a *DocBook* publishing system.

1.2. Pull-quote Conventions

Terminal output and source code listings are set off visually from the surrounding text.

Output sent to a terminal is set in **mono-spaced roman** and presented thus:

```
books      Desktop  documentation  drafts  mss    photos  stuff  svn
books_tests Desktop1  downloads      images  notes  scripts svgs
```

Source-code listings are also set in **mono-spaced roman** but add syntax highlighting as follows:

```
package org.jboss.book.jca.ex1;

import javax.naming.InitialContext;

public class ExClient
{
    public static void main(String args[])
        throws Exception
    {
        InitialContext iniCtx = new InitialContext();
        Object          ref    = iniCtx.lookup("EchoBean");
        EchoHome        home   = (EchoHome) ref;
        Echo             echo   = home.create();

        System.out.println("Created Echo");

        System.out.println("Echo.echo('Hello') = " + echo.echo("Hello"));
    }
}
```

1.3. Notes and Warnings

Finally, we use three visual styles to draw attention to information that might otherwise be overlooked.



Note

Notes are tips, shortcuts or alternative approaches to the task at hand. Ignoring a note should have no negative consequences, but you might miss out on a trick that makes your life easier.



Important

Important boxes detail things that are easily missed: configuration changes that only apply to the current session, or services that need restarting before an update will apply. Ignoring a box labeled 'Important' will not cause data loss but may cause irritation and frustration.



Warning

Warnings should not be ignored. Ignoring warnings will most likely cause data loss.

2. Getting Help

2.1. Do You Need Help?

If you experience difficulty with a procedure or other information described in this documentation, visit the Red Hat Knowledgebase at <http://kbase.redhat.com> to search or browse through technical support articles about Red Hat products, or visit the Red Hat Customer Portal at <http://access.redhat.com>. You can also access the OpenShift web site at <https://openshift.redhat.com/> to find blogs, FAQs, forums, and other sources of information.

Red Hat also hosts a large number of electronic mailing lists for discussion of Red Hat software and technology. You can find a list of publicly available mailing lists at <https://www.redhat.com/mailman/listinfo>. Click the name of any mailing list to subscribe to that list or to access the list archives.

2.2. We Need Feedback!

If you find a typographical or any other error in this manual, or if you have thought of a way to make this manual better, we would love to hear from you! Please submit a report in Bugzilla: <http://bugzilla.redhat.com/> against the product **OpenShift**.

When submitting a bug report, be sure to mention the manual's identifier: *Docs User Guide*

If you have a suggestion for improving the documentation, try to be as specific as possible when describing it. If you have found an error, please include the section number and some of the surrounding text so we can find it easily.

OpenShift Architecture Overview

1.1. Platform Overview

OpenShift enables you to create, deploy and manage applications within the cloud. It provides disk space, CPU resources, memory, network connectivity, and an Apache or JBoss server. Depending on the type of application you are building, you also have access to a template file system layout for that type (for example, PHP, Python, and Ruby/Rails). OpenShift also generates a limited DNS for you.

OpenShift provides dedicated `/var/tmp` and `/tmp` directories for each user application. The `/var/tmp` directory is a symbolic link to `/tmp`. Each `/tmp` directory is completely isolated from the `/tmp` directories of all other applications. Files not touched in these directories for any 10-day period are deleted.

The two basic functional units of OpenShift are the Broker, which provides the interface, and the Cartridges, which provide application frameworks.

Broker

The Broker is the single point of contact for all application management activities. It is responsible for managing user logins, DNS, application state, and general orchestration of the application. Customer interaction with the broker is generally performed using either the Web console, CLI tools, or JBoss tools, using a REST-based API.

Cartridges

Cartridges provide the actual functionality necessary to run applications. Numerous cartridges are currently available to support languages such as Perl, PHP, and Ruby, as well as many database cartridges, such as PostgreSQL, MySQL, and MongoDB.




Figure 1.1. High-level OpenShift Platform Overview

1.2. System Resources and Application Containers

The system resources and security containers provided by the OpenShift platform are gears, nodes and districts.

- Gears — Gears provide a resource-constrained container where you can run one or more cartridges. They limit the amount of RAM and disk space available to a cartridge.



Figure 1.2. Application Cartridges on Gears

OpenShift currently provides two gear sizes:

- Small — provides 512MB of RAM, 100MB of swap space, and 1GB of disk space
- Medium — provides 1GB of RAM, 100MB of swap space, and 1GB of disk space

By default, you can use up to three small gears (a total of 1.5GB of RAM and 3GB of disk space). OpenShift can assign these three gears to a single application and its cartridges (Cron, MySQL, etc.) or you can use each gear for a separate application.

- Nodes — To enable resource sharing, multiple gears run on a single physical or virtual machine. This machine is referred to as an OpenShift node. Gears are generally over-allocated on nodes because not all applications are active at the same time.
- Districts — Districts define a set of nodes within which gears can be easily moved to load-balance the resource usage of nodes. This means that no node gets overloaded with many, heavily-used gears.

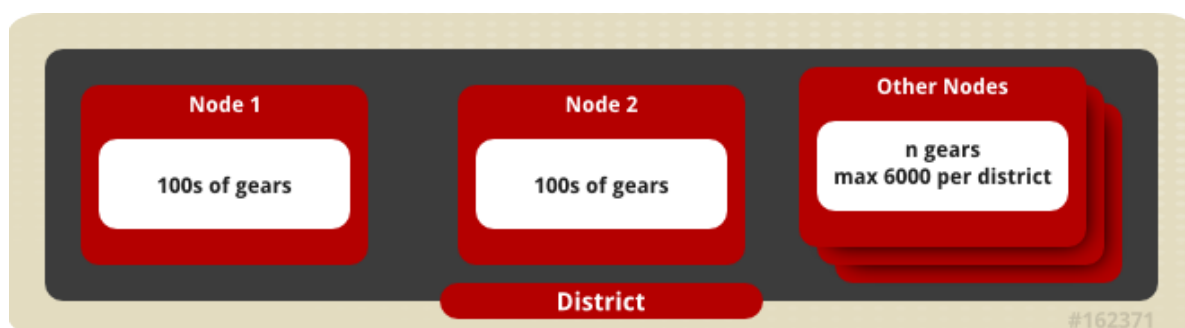


Figure 1.3. Nodes in a District

1.3. OpenShift Applications

Each username can only support a single namespace, but you can create multiple applications within this namespace. If you need multiple namespaces, you need to create multiple accounts using different usernames.

The following diagram illustrates the relationships between usernames, namespaces, domain names, and applications:



Figure 1.4. Name Relationships

Applications consist of a number of components:

- **Namespace** — The namespace is not directly related to DNS, instead it provides a unique group identifier for all the applications of a specific user. The namespace is appended to the application name to form a final application URL of the form `http://[APP NAME]-[NAMESPACE].rhcloud.com`
- **Application Name** — The user-selected name of the application. The final URL to access the application is of the form `http://[APP NAME]-[NAMESPACE].rhcloud.com`
- **Aliases** — Users can provide their own DNS names for the application by registering an *alias* with OpenShift and pointing the DNS entry to the OpenShift servers.
- **Application GIT repository** — A GIT repository is created for each application. You can modify your code in the repository and then run **git push** to deploy the revised code.
- **Application dependencies** — You need to specify the cartridges required to run your application. The following table describes the functions provided by cartridges:

Table 1.1. Cartridge Functions

Function	Description
Framework (required)	These cartridges allow you to choose from a variety of programming languages and frameworks for developing your application. Every application requires a framework cartridge. Examples include PHP, JBoss, and Ruby.
Database	These cartridges provide your application with one of several database back ends. Examples include MySQL, MongoDB, and PostgreSQL.
Database management	These cartridges provide functionality for managing your application's database using third-party software. Examples include phpMyAdmin, RockMongo, and 10gen's MongoDB Monitoring Service (MMS).
Monitoring and Management	These cartridges provide a range of options for managing and monitoring your application. Examples include the Cron task scheduler, the Jenkins Client, and OpenShift Metrics.

OpenShift supports the use of **package.json** as a mechanism for applications to specify dependencies and other requirements. When you deploy or update your application, OpenShift invokes **npm**, the NodeJS package manager, to process that file. Refer to <http://npmjs.org/doc/json.html> for information about the NodeJS package manager.

1.3.1. Using Scaled Applications

Application scaling provides for the automatic allocation of resources based on demand. OpenShift monitors the resource requirements of scaled applications, and increases or decreases available resources accordingly. You need to specify whether an application is scaled or not when you create the application. A scaled application cannot be converted to a non-scaled application. Likewise, a non-scaled application cannot be converted to a scaled application.

Scaled versus Non-scaled Applications

If you create a non-scaled application, it only consumes one of the default quota of gears assigned to users. That is, it only consumes one of the available three gears. If you create a scaled application, it consumes two of the available gears; one for the high-availability proxy (HAProxy) itself, and one for your actual application. If you add MySQL to your application, it is installed in its own dedicated gear.

How Scaling Works

Each application created on OpenShift always has the web cartridge associated with it. The web cartridge can, for example, be a PHP cartridge. When an application is scaled, a second cartridge, called HAProxy, is added to the application. The HAProxy cartridge listens to all incoming web page requests for an application and passes them on to the web cartridge, following defined guidelines for load monitoring.

As the number of web page requests to an application increase, the HAProxy will inform OpenShift when an overload of requests is detected. OpenShift will then create a copy of the existing web cartridge on a separate gear. In such a case, the web cartridge now has been scaled up two times. This process is repeated as more web page requests are detected by the HAProxy cartridge, and each time a copy of the web cartridge is created on a separate gear, the application scale factor increases by one.

However, not all OpenShift applications can be scaled, as detailed in the table below.

Table 1.2. Applications that can or cannot be scaled

Type of Application	Scalable
JBoss Application Server 7.1	Yes
JBoss Enterprise Application Platform 6.0	Yes
PHP 5.3	Yes
Python 2.6	Yes
Perl 5.10	Yes
Ruby 1.9	Yes
Ruby 1.8.7	Yes
Node.js 0.6	Yes
Jenkins	No
HAProxy	No
DIY	No

**Note**

You can only add MySQL 5.1, Jenkins Client 1.4, or MongoDB 2.0 cartridges to scaled applications.

Scaled applications can be created with the OpenShift client tools using the CLI.

1.4. Cartridges

Cartridges are the components of an OpenShift application, and include databases, build systems and management capabilities. Adding a cartridge to an application provides the desired capability without requiring you to administer or update the included feature.

Examples of cartridges include the different language cartridges (PHP, Ruby, etc.) that you select from when creating an application, database cartridges such as PostgreSQL, MySQL, and MongoDB, and management cartridges such as the cron scheduler cartridge that lets you set up timed jobs for your applications.

At the time of writing, the following cartridges are available:

Database Cartridges

- MySQL Database 5.1 — MySQL is a multi-user, multi-threaded SQL database server
- MongoDB NoSQL Database 2.0 — MongoDB is a scalable, high-performance, open source NoSQL database
- PostgreSQL Database 8.4 — PostgreSQL is an advanced Object-Relational database management system

Management Cartridges

- phpMyAdmin 3.4 — phpMyAdmin is a web-based MySQL administration tool
- RockMongo 1.1 — RockMongo is a web-based MongoDB administration tool
- 10gen MMS agent 0.1 — 10gen's MongoDB Monitoring Service (MMS)
- Jenkins Client 1.4 — a client for managing Jenkins-enabled applications
- HAProxy 1.4 — a high-performance TCP/HTTP load balancer
- Cron 1.4 — Cron is a daemon that runs specified programs at scheduled times
- OpenShift Metrics 0.1 — OpenShift Metrics is an experimental cartridge for monitoring applications

To add cartridges using the web Management Console, refer to [Section 2.7.1, “Adding Cartridges”](#). To add cartridges using the command line client tools, refer to [Section 3.6.1, “Adding Cartridges with the CLI”](#).

OpenShift Web Interface



Note

Before you start working through the procedures and examples in the following chapters, refer to the *Getting Started Guide* to ensure that you have performed all of the necessary steps to set up your environment for OpenShift.

2.1. Accessing the OpenShift Management Console

OpenShift applications can be created and managed using the *OpenShift Management Console*, a graphical user interface accessed with a browser. The Management Console also allows you to manage your OpenShift account settings and provides links to OpenShift documentation and community resources.

Procedure 2.1. To access the OpenShift Management Console:

1. On the OpenShift homepage, click **SIGN IN TO MANAGE YOUR APPS** in the upper-right corner.
2. On the sign in screen, enter your login name and password details then click **Sign in**.

Sign in to OpenShift

Login

Password

[Forgot your password?](#)

or [create an account](#)

Figure 2.1. Sign In Screen

If you do not have any applications, the **Create a New Application** screen opens. If you have applications, you are taken to the **My Applications** screen.

2.2. Managing Your OpenShift Account

2.2.1. Changing Your Password

Your OpenShift user account password can be changed using the Management Console.

When choosing a password we recommend using a combination of numbers, symbols, and upper and lower case letters for extra security.

Procedure 2.2. To change your password:

1. Access the Management Console and click **My Account** in the navigation bar at the top of the page.
2. In the **Personal Information** section, click **Change password** then follow the on-screen instructions.

2.2.2. Resetting Your Password

If you forget your OpenShift user account password, you can have a new password sent to your e-mail address.

Procedure 2.3. To reset your password:

1. On the OpenShift homepage, click **SIGN IN TO MANAGE YOUR APPS** in the upper-right corner.
2. On the sign in screen, click **Forgot your password?**
3. Enter your e-mail address and click **Reset Password**.

A new password is sent to your e-mail address. Use this password to access the OpenShift Management Console and change your password.

2.2.3. Changing Namespaces

Changing your namespace deletes the old namespace and creates a new one. It also automatically updates the public URLs and repository addresses for your applications. In order to **git push** future changes to your applications after changing your namespace, the git **config** file must be updated with the new repository address.

OpenShift uses a blacklist to restrict the list of available namespace and application names that you can use. This list is maintained on the server. If you try to change your namespace to any members of this blacklist, the command will fail.




Important

This procedure alters the URLs for your applications. You need to update any bookmarks and external links you have made to these URLs. Links made using an alias do not need to be changed. See [Section 3.4.2.2, “Using Arbitrary DNS Names”](#).

Procedure 2.4. To change your namespace:

1. Access the Management Console and click **My Account** in the navigation bar at the top of the page.
2. Scroll down to the **Namespace** section and click **Change your namespace**.
3. Enter your desired namespace in the box provided and click **Save**.



Change your namespace

Warning Changing your namespace will alter the public URLs of your applications. If you change your namespace you will need to update any bookmarks or external links to your applications.

http://applicationname-- | .rhcloud.com

Your namespace must be letters or numbers with no spaces or symbols.


Cancel Save

Figure 2.2. Change Namespace Dialog

The public URLs and repository addresses of your applications automatically update with the new namespace. To enable the **git push** command to function properly, update the git **config** file using the following procedure.

Procedure 2.5. To update the git config files:

1. Access the Management console and click **My Applications** in the navigation bar at the top of the page.
2. Click on your first application.
3. Copy the entire SSH address located in the **GIT REPOSITORY** box.
4. Open the git **config** file located in *path/to/appdirectory/.git/* and replace the remote URL address with the new SSH address for your application.
5. Repeat the previous four steps for each of your applications to update their git **config** files.



crossword Delete

Cartridges Add Cartridge +

PHP 5.3 CREATED ABOUT 1 MONTH AGO

GIT REPOSITORY

ssh://ad92e6d82188429b8218b7511f164e25@crossword-gametime.rhcloud.com/~/.git/crossword.git/

ADD CARTRIDGE

VISIBILITY
External

URL
http://crossword-gametime.rhcloud.com

ALIASES
No alias set

CLIENT TOOLS HELP
OpenShift User Guide
Sync your OpenShift repo with an existing Git repo

Figure 2.3. Application Details Screen Showing Git Repository Address

2.2.4. Managing SSH Keys

Using the OpenShift Management Console you can add, remove, and update public keys to control the access of other contributors to your OpenShift applications.

To view the public keys associated with your account, access the Management Console and click **My Account** in the navigation bar at the top of the page. Your current SSH public keys are listed in the **Public Keys** section.

2.2.4.1. Generating New Keys

The **ssh-keygen** command generates a new pair of RSA or DSA keys as specified. You can then use the Management Console to add the new keys to your account.

Procedure 2.6. To generate new SSH keys with the **ssh-keygen** command:

1. Manually generate a new pair of keys, replacing *KeyType* with the type of key you want to generate, either *dsa* or *rsa*:

```
$ ssh-keygen -t KeyType
```

2. By default the new SSH keys are located in the `/home/username/.ssh/` directory.

2.2.4.2. Adding a Key

Procedure 2.7. To add a key:

1. Access the Management Console and click **My Account** in the navigation bar at the top of the page.
2. In the **Public Keys** section, click **Add a new key**.
3. Enter a name for your key then paste the public key in the space provided.
4. Click **Create** to add your public key.



Important

If you copy and paste your SSH key from an editor or terminal with the word wrap function enabled, the key may include unnecessary line breaks. Therefore, the OpenShift web console will reject the SSH key and the upload process will fail. Make sure that when you paste your key into the web console, the key contents are correct and do not contain any unnecessary line breaks.

2.2.4.3. Removing a Key

Procedure 2.8. To remove a key:

1. Access the Management Console and click **My Account** in the navigation bar at the top of the page.
2. In the **Public Keys** section, click **Delete** next to the key you want to remove.
3. A dialog box appears asking you to confirm the deletion. Click **OK** to confirm.

2.3. Creating Applications

Creating applications using the OpenShift Management Console is a simple process.

2.3.1. Creating Basic Applications

OpenShift provide a selection of application types that you can use to build your applications.

Procedure 2.9. To create an application:

1. Access the Management Console and click **Create a New Application** in the navigation bar at the top of the page.
2. Choose the application type you want to create from the **Web Cartridges** section and click **Select**. The application types currently available on the Management Console are:
 - PHP 5.3 — for PHP applications
 - Python 2.6 — for Web Server Gateway Interface applications
 - Ruby 1.9 — for Ruby Webserver Interface applications
 - Ruby 1.8.7 — for Ruby Webserver Interface applications
 - Perl 5.10 — for Perl applications
 - JBoss Application Server 7.1 — for JBoss AS applications running on Java 6 or Java 7
 - JBoss Enterprise Application Platform 6.0 — for JBoss EAP applications running on Java 6 or Java 7
 - Node.js 0.6 — for Node.js applications
 - Jenkins Server — a continuous integration (CI) server that enables complex builds
 - Do-It-Yourself (DIY) — a blank slate for trying unsupported languages, frameworks, and middleware on OpenShift
3. Type a name for your application in the box provided and click **Create Application**.

PHP 5.3

Website: <http://www.php.net>
Version: 5.3.2

PHP is a general-purpose server-side scripting language originally designed for Web development to produce dynamic Web pages. The mod_php Apache module is used to execute PHP applications. Popular development frameworks include: CakePHP, Symfony, and Code Igniter. Popular applications include: Drupal, Wordpress, and Mediawiki.

Public URL

Your application name uniquely identifies the application and becomes part of your public URL. You can add your own domain names to the application later.

Figure 2.4. Create Application Dialog

2.3.2. Creating Preconfigured Applications

In addition to the standard application web cartridges, OpenShift provides several preconfigured applications that allow you to create complete applications quickly and easily. These preconfigured applications are automatically created with a web cartridge, any other required cartridges (such as a database), and all of the necessary code for a fully functioning application.

Procedure 2.10. To create a preconfigured application:

1. Access the Management Console and click **Create a New Application** in the navigation bar at the top of the page.

2. Choose the preconfigured application you want to create from the **Instant Applications** section and click **Select**. The preconfigured applications currently available on the Management Console are:
 - CakePHP — an open source web application framework written in PHP
 - Drupal — an open source content management platform written in PHP
 - Kitchensink Example — a JBoss AS7 application that demonstrates the new features of Java EE 6
 - Ruby on Rails — an open source web application framework for Ruby
 - Spring Framework — an open source web application framework that runs on JBoss EAP 6.0
 - Wordpress — a semantic personal publishing platform written in PHP with a MySQL back end
3. Type a name for your application in the box provided and click **Create Application**.

As preconfigured applications are more complicated to build than basic applications, they can take longer to become available online. If you receive a 503 Service Unavailable error when attempting to view your preconfigured application after you create it, wait a few minutes and try again.

2.3.3. Cloning Application Files

After you have created an application using the OpenShift Management Console, run the **git clone** command to copy the application's remote repository into your local working directory.

Procedure 2.11. To clone the remote repository:

1. Click **My Applications** in the navigation bar at the top of the page and click on the application you want to clone.
2. Copy the entire SSH address located in the **GIT REPOSITORY** box.
3. Open a terminal and use the following command to clone the remote repository to the working directory, replacing the example SSH address with the address for your application:

```
$ git clone ssh://a261d0fc2932413694456e3473fdc972@crossword-gametime.rhcloud.com/~git/crossword.git/
```

The **git clone** command copies the template application files from the remote repository into the working directory. Edit the template application files to develop your own application.

2.4. Monitoring Application Resources

As described earlier in [Section 1.3.1, "Using Scaled Applications"](#), scaled applications are automatically allocated increased OpenShift resources based on demand. A scaled application can be allocated multiple gears as the demand increases. The amount of resources utilized by an application can be easily monitored and viewed from the OpenShift Management Console. Follow the instructions below to monitor and view application resources. The instructions below assume that a scaled application has already been created.

Procedure 2.12. To monitor application resources:

1. After you have created a scaled application, access the OpenShift Management Console.

- Click on the **My Applications** tab to view all of your applications, as highlighted in the figure below.

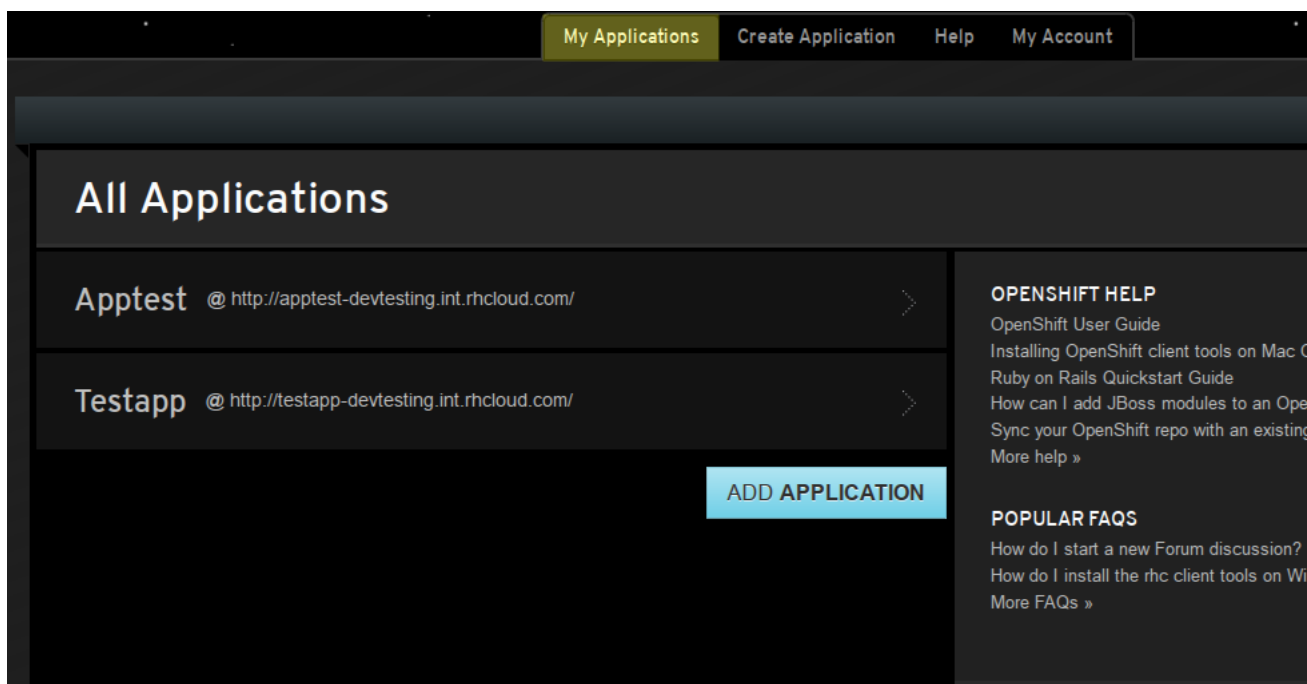


Figure 2.5. My Applications

- Under the **All Applications** section, click on the name of the scaled application for which you wish to monitor resource information, as highlighted in the figure below.

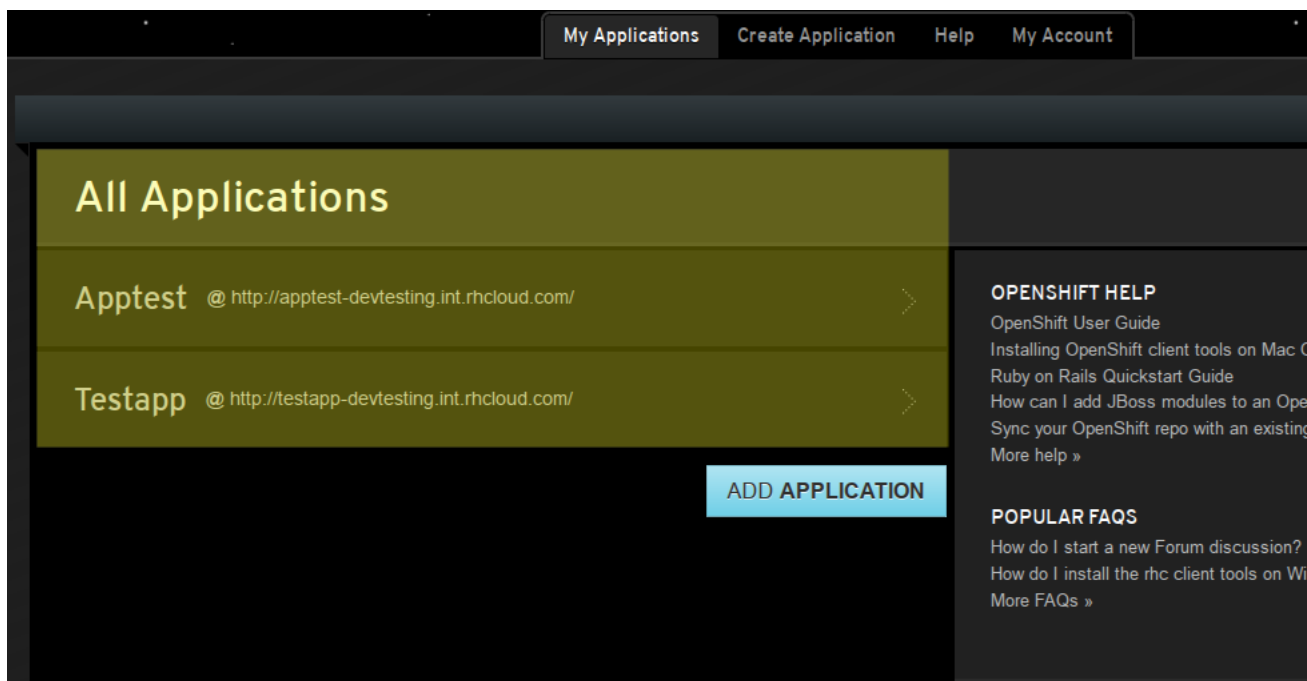


Figure 2.6. Select Application to Monitor

4. The OpenShift Management Console will display the number of gears, along with the size of the gears, used by the selected application, as highlighted in the figure below.

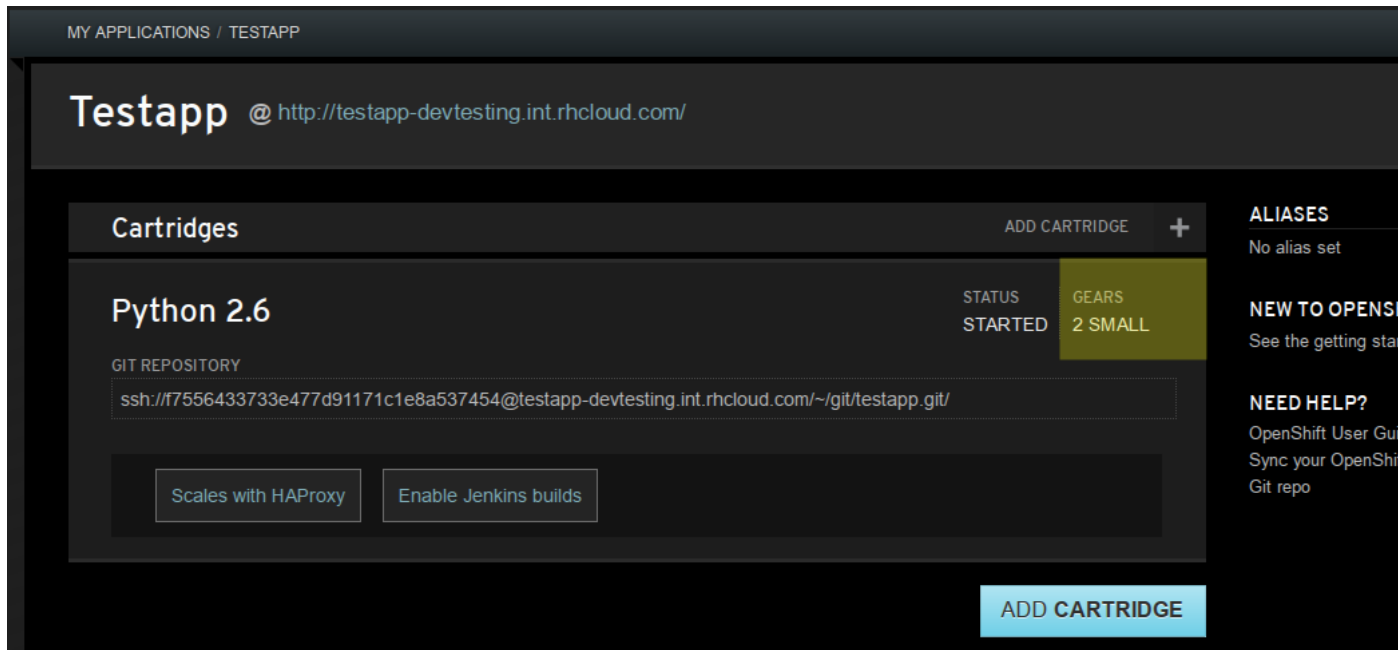


Figure 2.7. Application Resource Information

5. If you hover over the gear size information with your mouse, a popup message will display more detailed information, showing exactly how the gears are being utilized, as shown in the figure below.

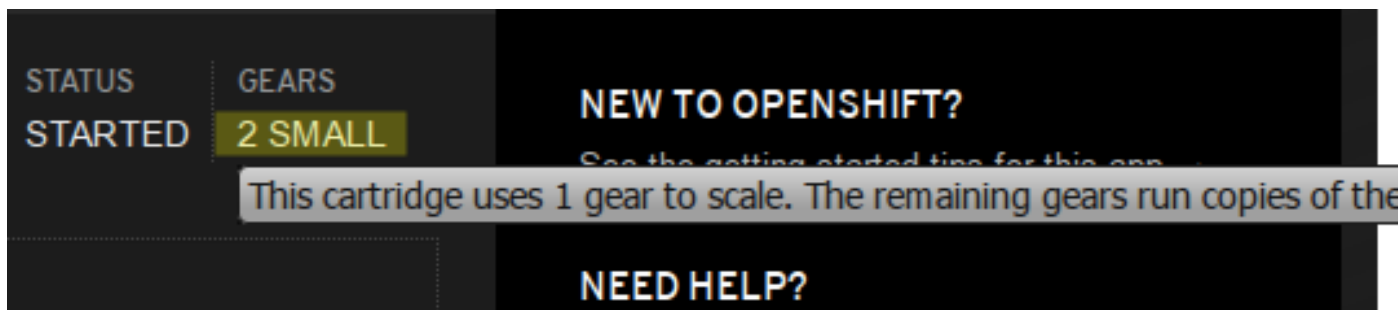


Figure 2.8. Application Gear Details

6. Because a scaled application is monitored by the HAProxy cartridge, this information is displayed in the OpenShift Management Console under your application information, as highlighted in the

figure below. Click **Scales with HAProxy** to get information about testing the scaling function of your application.

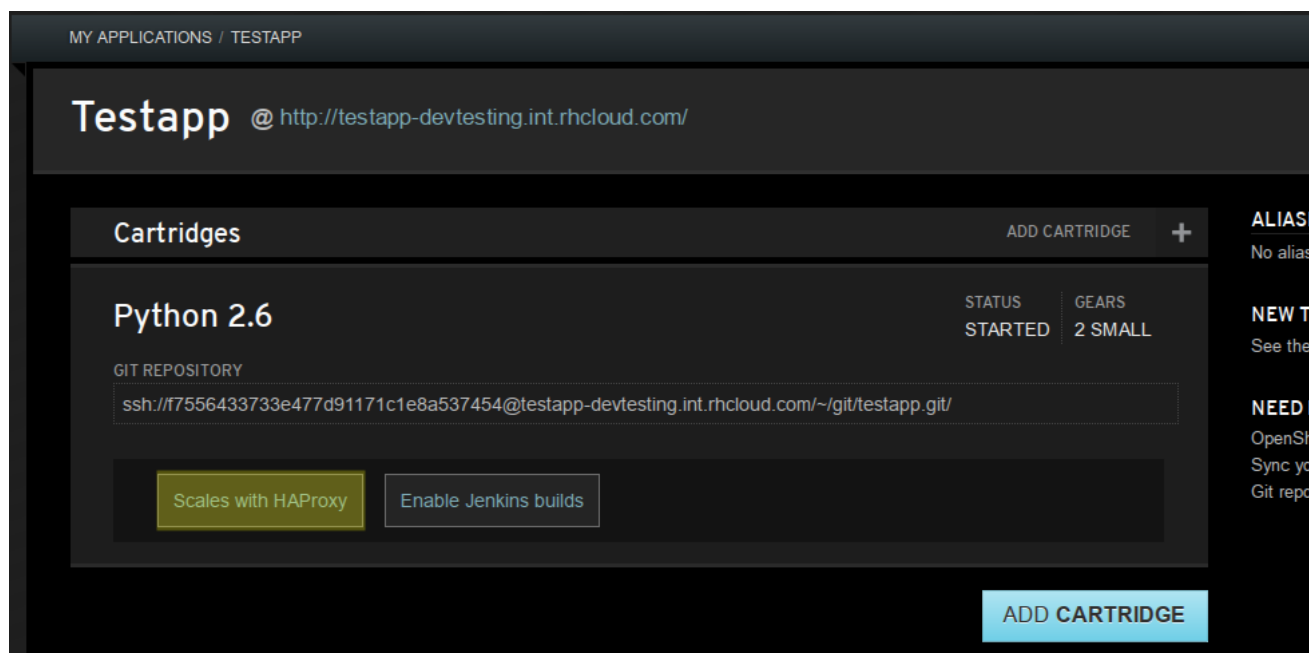


Figure 2.9. HAProxy Information



Note

At the time of this writing, the scaling function cannot be disabled from an application. The only way to disable scaling is to remove the scaled application, and create a new application without the scaling option.

2.5. Editing and Deploying Applications

To deploy your application to the cloud, you need to make any required changes to your application code base, commit those changes to your local repository, and then update the remote repository. Application files are stored in the local git repository that was cloned in the [Section 2.3, “Creating Applications”](#) procedure.

For advanced deployment options, refer to [Section 3.5, “Editing and Deploying Applications”](#) in the Command Line Interface chapter.

2.5.1. Preparing Your Application for Deployment

The `git clone` command used in [Section 2.3, “Creating Applications”](#) creates a starting point, or template, for you to create and develop your own applications. To synchronize your application files with the remote cloud repository, you need to commit all of your files to the appropriate directories in the local git repository, and then push them to the remote repository. For example, you may be developing your own PHP application and need to add new files and directories to the `$_ENV['OPENSIFT_APP_NAME']/php/` or other directories.

Procedure 2.13. To prepare your application for deployment:

1. Open a terminal and use the following command to add each new file or directory to the git index:

```
$ git add /path/to/newfile
```

2. Use the following command to commit your application files to your local repository, replacing *commit message* with your own message:

```
$ git commit -m "commit message"
```

Your application files are now ready to be deployed to the cloud.

2.5.2. Deploying Your Application to the Cloud

After you have committed your application files to the local repository, you need to push them to the remote repository. OpenShift will automatically stop, build, and restart your application with the committed changes.

Procedure 2.14. To deploy your application to the cloud:

- Use the following command to deploy your application to the remote repository:

```
$ git push
```

Your updated application is now publicly available.

2.6. Deleting Applications

If you no longer need a particular application, you can choose to delete it.

2.6.1. Deleting Remote Application Data



Warning

Application removal is irreversible. All remote data associated with the application will be removed.

Procedure 2.15. To delete an application:

1. Access the Management Console and click **My Applications** in the navigation bar at the top of the page.
2. Click on the application you want to delete.
3. Click **Delete** next to the application name.
4. You are asked to confirm the request. Click **Delete** to confirm.

This process deletes your remote application data. If you want to delete application data stored on your local machine, you must do so manually.

2.6.2. Deleting Local Application Data



Warning

The following procedure deletes the selected directory and all the files it contains. Ensure you enter the correct directory and that you no longer need the files it contains before running this command.

Procedure 2.16. To delete local application data

- Open a terminal and use the following command to delete the application data stored on your local machine:

```
$ rm -rf ~/path/to/app_directory/
```

2.7. Using Cartridges

The OpenShift Management Console provides an intuitive interface for managing your application's cartridges.

You can view the cartridges associated with an application by clicking **My Applications** in the navigation bar at the top of Management Console then clicking on the application you want to view.

2.7.1. Adding Cartridges

The OpenShift Management Console allows you to add cartridges to your applications easily. Note that certain cartridges are only available to be added after a prerequisite cartridge is added. Cartridge descriptions in the Management Console detail these dependencies.



Note

You can only add MySQL 5.1, Jenkins Client 1.4, or MongoDB cartridges to scaled applications.

Procedure 2.17. To add a cartridge:

1. Access the OpenShift Management Console and click **My Applications** in the navigation bar at the top of the page.
2. Click on the application to which you want to add a cartridge.
3. Click **Add Cartridge**.
4. Choose the cartridge to add to your application and click **Select**.
5. Click **Add Cartridge** to confirm the selection.

OpenShift Command Line Interface



Note

Before you start working through the procedures and examples in the following chapters, refer to the *Getting Started Guide* to ensure that you have performed all of the necessary steps to set up your environment for OpenShift.

3.1. Secure Shell Authentication

OpenShift uses the Secure Shell (SSH) network protocol to authenticate your account credentials to the OpenShift servers for secure communication. Successful authentication is necessary to manage your cloud environment, and OpenShift supports both RSA and DSA keys for SSH authentication. This section describes briefly how OpenShift authentication works, and provides information on how to manage SSH keys for OpenShift user accounts.

For successful authentication, the public SSH key on your computer must match the public key that has been uploaded to the OpenShift server. When you perform the initial configuration of the OpenShift client tools, the interactive setup wizard generates a new pair of SSH keys in the default `.ssh` folder of your home directory. The SSH key pair consists of the public key, `id_rsa.pub`, and the private key, `id_rsa`. As part of the initial configuration, you have the option of automatically uploading the public key, `id_rsa.pub`, to the OpenShift server. If you have not performed the initial configuration of the OpenShift client tools, refer to the *Getting Started Guide* for more information.

3.1.1. Resolving Authentication Issues

Occasionally your local public key may not match the public key for your account on the OpenShift server, or your key may not be found on the local file system. This can cause connection issues, or the SSH key authentication process can fail, in which case a new pair of SSH keys must be generated. If you are having problems authenticating, there are two ways you can generate a new pair of SSH keys:

- Use the interactive setup wizard (recommended)
- Generate and upload SSH keys manually

3.1.1.1. Using the Interactive Setup Wizard

The recommended method of resolving authentication issues is to use the interactive setup wizard to generate a new pair of SSH keys. The interactive setup wizard will also provide the option of automatically uploading your new public key to the OpenShift server. Use the command as shown below to launch the setup wizard and follow the onscreen instructions.

```
$ rhc setup
```

Refer to the *Getting Started Guide* for more information about the OpenShift client tools interactive setup wizard.

3.1.1.2. Generating and Uploading SSH Keys Manually

Although not recommended, advanced users can manually generate a new pair of SSH keys, and upload the public key to the OpenShift server. Follow the instructions below.

Procedure 3.1. To manually generate a new pair of SSH keys and upload the public key to the server:

1. Generate a new pair of keys:

```
$ ssh-keygen -t <KeyType>
```

where *KeyType* is the type of key you want to generate, either DSA or RSA.

2. Add the new public key to the user account:

```
$ rhc sshkey add -i <KeyName> -k <KeyPath> -l <UserName>
```

where *<KeyPath>* is the path and filename of the public key that you want to add, and *<KeyName>* is a name that you specify to identify this key.

3. Add the new public key to the SSH agent:

```
$ ssh-add <KeyPath>
```

3.1.2. Managing Multiple SSH Keys

OpenShift provides command line tools to manage multiple SSH keys to control access to applications among different users. Multiple SSH key management is provided with the **rhc sshkey** command, with options available to view, add, update, and remove public keys on the OpenShift server. Refer to the subsequent sections for more information.

3.1.2.1. Viewing Public Keys

Use the **rhc sshkey list** command to view all public keys that have been added to the OpenShift server for the specified user account:

```
$ rhc sshkey list -l <UserName>
Password:

SSH keys
=====
Name: default
Type: ssh-rsa
Key: AAAAB3NzaC1yc2EAAAABIwAAAQEAQT20LZBFrbzq
```

3.1.2.2. Adding a Key

Use the **rhc sshkey add** command to add a public key to the OpenShift server for the specified user account:

```
$ rhc sshkey add -i <KeyName> [-k <KeyPath>] -l <UserName>
Password:
Success
```

where *<KeyName>* is the user-specified identifier for the SSH key, and *<KeyPath>* is the path to an existing key (optional). If an existing key is not specified, a new key is generated and uploaded to the server.

Refer to the **rhc sshkey** manual page for more information.

3.1.2.3. Updating a Key

Use the **rhc sshkey update** command to update an existing public key for the specified user account:

```
$ rhc sshkey update -i <KeyName> [-k <KeyPath>] -l <UserName>
```

where *<KeyName>* is the user-specified identifier for the SSH key, and *<KeyPath>* is the path to an existing key (optional). If an existing key is not specified, a new key is generated and uploaded to the server.

Refer to the **rhc sshkey** manual page for more information.

3.1.2.4. Removing a Key

Use the **rhc sshkey remove** command to remove an existing public key from the OpenShift server for the specified user account:

```
$ rhc sshkey remove -i <KeyName> -l <UserName>
```



Important

The primary, or **default** key, can only be modified; it cannot be deleted.

3.2. Working With Domains

3.2.1. Creating a Domain

You must create a domain before you can create an application. OpenShift uses non-strict domain names (that is, there is no preceding period), and the domain name forms part of the application name. The syntax for the application name is *ApplicationName-DomainName.rhcloud.com*.

Each username can only support a single domain, but you can create multiple applications within the domain. If you need multiple domains, you need to create multiple accounts using different usernames.

Use the **rhc domain create** command as shown below to create a new domain. Be sure to replace the example parameters with those that suit your deployment:

```
$ rhc domain create -n <DomainName> -l <rhlogin> -p <password> [OptionalParameters]
```

- *DomainName* — specifies the domain that you want to create. This must contain a maximum of 16 alphanumeric characters.
- *rhlogin* — this can be your Red Hat Network login or the email address that you used to apply for your OpenShift account.

- *password* — the password that you used to apply for your OpenShift account.



Note

The **rhc domain create** command creates a local config file in `~/.openshift/express.conf` and updates it with your rhlogin. The rhc client tools use the login specified in the `~/.openshift/express.conf` file by default when you run further commands. If no default login is configured, or if you want to use a different login, include the **-l rhlogin** parameter as a command line argument.

Optional Parameters

Further optional parameters exist that you can pass to the **rhc domain create** command. Refer to the manual page for details.

Example 3.1. Creating a new domain

The following example demonstrates how to create a new domain. Unnecessary details have been removed for clarification purposes.

```
$ rhc domain create -n ecs -l jsmith@mail.com

Password: *****

OpenShift key found at /home/user/.ssh/id_rsa. Reusing...
Creation successful
```



Note

If no SSH keys are found in the **.ssh** folder of your home directory, the **rhc domain create** command generates a new pair of keys for you. You will then be prompted to upload the public SSH key to the OpenShift server. Refer to [Section 3.1, “Secure Shell Authentication”](#) for more information.



Note

OpenShift uses a *blacklist* to restrict the list of available domain and application names that you can use. This list is maintained on the server. If you try to use **rhc domain create** or **rhc app create** with any members of this blacklist, the command will fail.

3.2.2. Altering a Domain

If you need to modify the name of an existing domain, use the **rhc domain alter** command as shown below.

**Important**

The **rhc domain alter** command alters the URLs of your applications. You need to update any bookmarks and external links you have made to these URLs. Links made using an alias do not need to be changed. Refer to [Section 3.4.2.2, “Using Arbitrary DNS Names”](#) for more information.

```
$ rhc domain alter -n <DomainName>
Password: *****

OpenShift key found at /home/user/.ssh/id_rsa. Reusing...
Alteration successful.
```

**Note**

If no SSH keys are found in the **.ssh** folder of your home directory, the **rhc domain alter** command generates a new pair of keys for you. You will then be prompted to upload the public SSH key to the OpenShift server. Refer to [Section 3.1, “Secure Shell Authentication”](#) for more information.

3.2.3. Deleting a Domain

If you no longer need a particular domain, you can use the **rhc domain destroy** command to remove, or *destroy*, that domain. Before you can destroy a domain, however, you need to ensure that it does not contain any applications.

Procedure 3.2. How to destroy a domain

1. Ensure that the domain does not contain any applications.

```
$ rhc domain show
Password:
Application Info
=====
```

If any applications are listed, use the **rhc app destroy** command to remove them.

```
$ rhc app destroy -a <ApplicationName>
```

**Warning**

Application removal is irreversible. All remote data associated with the application will be removed.

2. Use the **rhc domain destroy** command to destroy the domain.

```
$ rhc domain destroy -n <DomainName>
Password:
Success
```

After you have destroyed your domain, you need to create a new one before you can create any new applications or use the other client tools.

3.3. Viewing User Information

Use the **rhc domain show** command to display information about all of the current applications. You can view the *Framework Type*, *Creation Date*, *GitURL* and *PublicURL* details for each application.

The following example demonstrates how to display application information for user:

```
$ rhc domain show
Password: <user password>
Contacting https://openshift.redhat.com
Application Info
=====
racer
Framework: php-5.3
Creation: 2011-04-13T08:35:02-04:00
Git URL:
ssh://99b3a77fb9204c6687618bcb92c0548c@racer-ecs.rhcloud.com/~/.ssh/
Public URL: http://racer-ecs.rhcloud.com/
```

3.4. Creating Applications

3.4.1. Introduction

OpenShift provides both browser-based and command-line tools for you to create and deploy your applications. If you use the command-line tools to create a new application, a new, remote git repository is created and automatically cloned to your current directory on your local machine. Further, the hostname and IP address of your application is added to the list of known hosts (**~/.ssh/known_hosts**).

Prerequisites

Application creation requires a reliable network connection. The **rhc app create** command only makes a single attempt to create your application. OpenShift makes up to seven checks for the DNS entry to exist, and returns a failure message if not found.

If you continue to experience timeout issues, you may need to use the **--timeout** option on the command line to override the default values. OpenShift uses two timeout parameters: a *connection* timeout, which determines how long the client tries to connect to the server before timing out; and a *read* timeout, which determines how long the client waits for a response from the server. The default connection timeout value is 20 seconds. The default read timeout value is 120 seconds.

The `--timeout` option affects both timeout parameters, but it can only be used to *increase* values from the default. You cannot set the timeout value to be less than the default. For example, if you use `--timeout 50`, it sets the connection timeout value to 50 seconds, but does not affect the read timeout value. If you use `--timeout 150`, it sets both the connection and read timeout values to 150 seconds.

3.4.2. Creating Applications

Use the `rhc app create` command to create OpenShift applications. This command creates an application with the specified name, and in your domain namespace.

The following example demonstrates creating a PHP application called "racer" in the ecs domain that was created in [Example 3.1, "Creating a new domain"](#):

```
$ rhc app create -a racer -t php-5.3
Password: <user password>

Creating remote application space: racer
Contacting https://openshift.redhat.com

RESULT:
Successfully created application: racer

Adding rhcloud.com to ~/.ssh/config
Warning: Permanently added 'racer-ecs.rhcloud.com,50.17.130.104' (RSA) to the
list of known hosts.
Receiving objects: 100% (19/19), done.
Confirming application racer is available
Attempt # 1

Success! Your application is now published here:

http://racer-ecs.rhcloud.com/

The remote repository is located here:

ssh://99b3a77fb9204c6687618bcb92c0548c@racer-ecs.rhcloud.com/~/.git/racer.git/
```

As mentioned in [Section 3.2.1, "Creating a Domain"](#), each domain can support multiple applications. This means that you could run the same command, change only the application name (for example, create an application named "monitor"), resulting in a domain containing the following applications:

- `http://racer-ecs.rhcloud.com`
- `http://monitor-ecs.rhcloud.com`

3.4.2.1. Creating a Scaled Application

To create a scaled application, use the `rhc app create` command with the `-s` (or `--scaling`) option, as shown below:

```
$ rhc app create -a appName -t appType -s
```

3.4.2.2. Using Arbitrary DNS Names

You can specify meaningful DNS names for your **OpenShift** applications so that you can use your own DNS entries instead of `rhcloud.com`.

For example, for the two applications in the previous section, you could create aliases that better suited your own domain and application purposes. You can use the **`rhc app add-alias`** command to create these aliases, as follows:

```
$ rhc app add-alias -a racer --alias "racer.indy.com"
```

Both `http://racer-ecs.rhcloud.com` and `http://racer.indy.com` point to the same DNS and display the same site.

If at any time you need to remove the alias, use the **`remove-alias`** command:

```
$ rhc app remove-alias -a racer --alias "racer.indy.com"
```

Command Options

The **`rhc app create`** command supports the following options:

- `-l | --rhlogin` — The login name of the user creating the application.
- `-p | --password` — The login password of the user creating the application. If not specified on the command line, the command prompts for this information.
- `-a | --app` — The name of the application to create.
- `-c | --cartridge` — The added cartridge to manage (required for the cartridge command).
- `-r | --repo` — The path to the local repository that you want to create. This repository must be empty. OpenShift currently only supports git repositories.
- `-n | --nogit` — Only create a remote repository; do not automatically create a local copy.
- `-t | --type` — The framework type to create. OpenShift currently supports the following application types:
 - `php-5.3` — for PHP applications
 - `python-2.6` — for Web Server Gateway Interface applications
 - `ruby-1.9` — for Ruby Webserver Interface applications
 - `ruby-1.8` — for Ruby Webserver Interface applications
 - `perl-5.10` — for Perl applications
 - `jbossas-7` — for JBoss AS applications running on Java 6 or Java 7
 - `jbosseap-6.0` — for JBoss EAP applications running on Java 6 or Java 7
 - `jenkins-1.4` — for Jenkins applications
 - `nodejs-0.6` — for Node.js applications
 - `diy-0.1` — a "Do it Yourself (DIY)" cartridge type used to create applications of no specific type

**Note**

You need to specify the version of the application type. This is to enable support for multiple versions in later releases of OpenShift.

- `-d` | `--debug` — Enable debugging mode to produce more verbose output to assist in troubleshooting.
- `-h` | `--help` — Display command usage information.

Refer to the **rhc app create** manual page for a complete list of options.

3.4.2.2.1. Using DIY Cartridges

You can use OpenShift to create applications of no specific type, for build or other testing purposes. The syntax for using this cartridge type is the same as all other cartridge types:

```
$ rhc app create -a myapp -t diy-0.1
```

This creates an application that is not publicly available nor does it have anything running. You need to use **git push** and the `.openshift/action_hooks/` scripts to perform all required operations.

3.5. Editing and Deploying Applications

To deploy your application to the cloud, you need to make any required changes to your application code base, commit those changes to your local repository, and then update the remote repository. Application files are stored in the local git repository that was cloned as part of the application creation process.

3.5.1. Preparing Your Application for Deployment

The **rhc app create** command creates a starting point, or template, for you to create and develop your own applications. To synchronize your application files with the remote cloud repository, you need to commit all of your files to the appropriate directories in the local git repository, and then push them to the remote repository. For example, you may be developing your own PHP application and need to add new files and directories to the `$_ENV['OPENSIFT_APP_NAME']/php/` or other directories.

Procedure 3.3. To prepare your application for deployment using the command line:

1. Use the following command to add each new file and directory to the git index:

```
$ git add /path/to/newfile
```

2. Use the following command to commit your application to your local repository:

```
$ git commit -m "commit message"
```

3.5.2. Deploying Your Application to the Cloud

After you have added your application files to the local repository, you need to push them to the remote repository. OpenShift will automatically stop, build, and restart your application with the committed changes.

Procedure 3.4. To deploy your application to the cloud using the command line:

- Use the following command to deploy your application to the remote repository:

```
$ git push
```

3.5.3. Hot Deploying Applications

When you run the **git push** command to upload code modifications, OpenShift stops, builds, deploys and restarts your application. This entire process takes time to complete and is unnecessary for many types of code changes. Hot deploying your application allows your changes to take effect without restarting the application cartridge, thus increasing deployment speed and minimizing application downtime.

OpenShift provides support for hot deployment through a **hot_deploy** marker file. If the marker is present, supported application cartridges automatically hot deploy when you run the **git push** command.

Table 3.1. Application types that can or cannot be hot deployed

Type of Application	Hot Deploy
JBoss Application Server 7.1	Yes
JBoss Enterprise Application Platform 6.0	Yes
PHP 5.3	Yes
Perl 5.10	Yes
Python 2.6	No
Ruby 1.9	No
Ruby 1.8.7	No
Node.js 0.6	No
Jenkins	No
HAProxy	No
DIY	No

3.5.3.1. Hot Deployment Build Details

JBoss AS 7.1 and JBoss EAP 6.0

When you hot deploy JBoss AS 7.1 and JBoss EAP 6.0 applications, the Maven build is executed (either with Jenkins or without), and during the deploy step the **app-root/repo/deployments** directory is synced with the JBoss deployments directory. The server does not restart. For example, if the sync results in a WAR being modified, the JBoss HDScanner notices the modification and re-deploys it. If it notices a WAR has been removed, it undeploys it.

PHP 5.3 and Perl 5.10

When you hot deploy PHP 5.3 and Perl 5.10 applications, the application code is built (dependencies are processed and user build action_hooks are run) and deployed to the application server. The server does not restart. This is true for both Jenkins and non-Jenkins enabled applications. For Jenkins enabled applications, the build is performed on a Jenkins slave instance and then synced to the gear(s) where the application server is running.

3.5.3.2. Enabling and Disabling Hot Deployment

Procedure 3.5. To enable hot deployment:

- Open a terminal and run the command applicable to your operating system from your application's root directory to create the **hot_deploy** marker file:

Windows

```
C:\app_directory> copy NUL > .openshift\markers\hot_deploy
```

Linux and Mac

```
[user@user app_directory]$ touch .openshift/markers/hot_deploy
```

Procedure 3.6. To disable hot deployment:

- Hot deployment is disabled by deleting the **hot_deploy** marker file. To delete this file run the command applicable to your operating system from your application's root directory:

Windows

```
C:\app_directory> del .openshift\markers\hot_deploy
```

Linux and Mac

```
[user@user app_directory]$ rm .openshift/markers/hot_deploy
```

3.5.4. Deploying JBoss Applications

3.5.4.1. Deploying on Java 6 or Java 7

You can run JBoss AS 7 and JBoss EAP 6.0 applications on either Java 6 or Java 7. The Java version is specified by a **java7** marker file in your application's **markers** directory. By default, new JBoss applications have the **java7** marker enabled. Applications without the **java7** marker are deployed using Java 6.

Procedure 3.7. To use Java 7:

- Open a terminal and run the commands applicable to your operating system from your application's root directory to create the **java7** marker file and update the remote repository:

Windows

```
C:\app_directory> copy NUL > .openshift\markers\java7
C:\app_directory> git add .openshift\markers\java7
C:\app_directory> git commit -a -m "Add Java 7 marker"
C:\app_directory> git push
```

Linux and Mac

```
[user@user app_directory]$ touch .openshift/markers/java7
[user@user app_directory]$ git add .openshift/markers/java7
[user@user app_directory]$ git commit -a -m "Add Java 7 marker"
[user@user app_directory]$ git push
```

Procedure 3.8. To use Java 6:

- Java 6 is enabled by deleting the **java7** marker file. Run the commands applicable to your operating system from your application's root directory to delete the **java7** marker file and update the remote repository:

Windows

```
C:\app_directory> del .openshift\markers\java7
C:\app_directory> git commit -a -m "Remove Java 7 marker"
C:\app_directory> git push
```

Linux and Mac

```
[user@user app_directory]$ rm .openshift/markers/java7
[user@user app_directory]$ git commit -a -m "Remove Java 7 marker"
[user@user app_directory]$ git push
```

3.5.4.2. Automatic Deployment

To automatically deploy your applications to the OpenShift server runtime, you can place your deployment content, for example, `.war`, `.ear`, `.jar`, and `.sar` files, in the JBoss Application Server (AS) or JBoss Enterprise Application Platform (EAP) **deployments/** directory.

The file system deployment scanner in JBoss AS 7 and JBoss EAP 6.0 relies on a system of marker files. You can add or remove various marker files and the scanner deploys, withdraws, or redeploys content based on the type of marker file.

These marker files always have the same name as the deployment content to which they relate, but with an additional file suffix to indicate the type of action to perform. For example, the marker file to indicate that the **example.war** file should be deployed is named **example.war.dodeploy**.

Option 1: Uploading content in a Maven src structure

You can upload your content in a Maven src structure similar to the example **ROOT.war** file. When you perform a **git push** the application is built and deployed. This is the typical deployment option, and requires that your **pom.xml** be stored at the root of your repository, and that you have a *maven-war-*

plugin similar to the one in the example file to move the output from the build to the **deployments/** directory.

By default the *warName* is **ROOT** within the **pom.xml** file. This will render the webapp contents at `http://app_name-namespace.rhcloud.com`. If you change the *warName* in **pom.xml** to **app_name**, then your base URL would become `http://app_name-namespace.rhcloud.com/app_name`.



Important

If you are building locally, add any output **.war** and **.ear** files in the **deployments/** directory from the build to your **.gitignore** file.

Option 2: Uploading prebuilt content

You can use **git push** to push prebuilt **.war** files (with the corresponding **.dodeploy** file for exploded **.war** files) into the **deployments/** directory.



Important

If you use this method using the default repository, first run **git rm -r src/ pom.xml** from the root directory of your repository. If the **pom.xml** file still exists, the build will run again and replace any prebuilt content.

3.5.4.3. Types of Marker Files

Different marker file suffixes have different meanings. The relevant marker file types are as follows:

Table 3.2. Sample Table

Marker File Type	Description
.dodeploy	Placed by the user to indicate that the given content should be deployed into the runtime (or redeployed if already deployed in the runtime.)
.deploying	Placed by the deployment scanner service to indicate that it has detected a .dodeploy file and is in the process of deploying the content. This marker file is deleted when the deployment process completes.
.deployed	Placed by the deployment scanner service to indicate that the given content has been deployed to the runtime. If you delete this file, the content will be withdrawn.
.faileddeploy	Placed by the deployment scanner service to indicate that the given content failed to deploy to the runtime. The content of this file will include some information about the cause of the failure.
.undeploying	Placed by the deployment scanner service to indicate that it has detected a .deployed file has been deleted and the content is being withdrawn. This marker file is deleted when the withdrawal process completes.
.undeployed	Placed by the deployment scanner service to indicate that the given content has been withdrawn from the runtime. If you delete this file, it has no impact.

3.5.4.4. Example JBoss Application Deployment Workflows

The following sections describe the basic workflows involved in deploying, withdrawing, and redeploying JBoss prebuilt applications to OpenShift.

To add new, zipped content and deploy it, run the following command:

```
$ cp target/example.war deployments/
```

To add new, unzipped content and deploy it, run the following commands:

```
$ cp -r target/example.war/ deployments/  
$ touch deployments/example.war.dodeploy
```

To withdraw deployed content, run the following command:

```
$ git rm deployments/example.war.dodeploy deployments/example.war
```

To replace currently deployed, zipped content with a new version and deploy it, run the following command:

```
$ cp target/example.war deployments/
```

To replace currently deployed, unzipped content with a new version and deploy it, run the following commands:

```
$ git rm -rf deployments/example.war/  
$ cp -r target/example.war/ deployments/  
$ touch deployments/example.war.dodeploy
```

3.6. Using Cartridges with the CLI

The OpenShift client tools provide a full range of commands for managing cartridges that you add to your applications.

You can view the cartridges associated with your applications by running the **rhc domain show** command. To view an individual application's information, use **rhc app show -a *AppName***. Installed cartridges are listed under the **Embedded** heading.

3.6.1. Adding Cartridges with the CLI

The current list of available cartridges can be viewed by running the **rhc app cartridge list** command. Note that certain cartridges can only be added after a prerequisite cartridge is added. Cartridge descriptions in the OpenShift website Management Console detail these dependencies.

**Note**

You can only add MySQL 5.1, Jenkins Client 1.4, or MongoDB cartridges to scaled applications.

Procedure 3.9. To add a cartridge using the CLI:

- Run the following command, replacing *AppName* and *CartType* with the name of the application to which you want to add a cartridge and the type of cartridge you want to add:

```
$ rhc app cartridge add -a AppName -c CartType
```

3.6.2. Managing Cartridges

You can use the following options with the **rhc app cartridge** command, specifying the application and cartridge to manage with the **-a AppName** and **-c CartType** arguments:

Table 3.3. Application management command argument options

Option	Details
list	Lists supported cartridges.
add	Adds a cartridge.
remove	Removes an added cartridge.
stop	Stops an added cartridge.
start	Starts an added cartridge.
restart	Restarts an added cartridge.
status	Returns the status of an added cartridge.
reload	Reloads the configuration of an added cartridge.

3.6.3. Working With Database Cartridges

OpenShift provides support for several database back ends, including PostgreSQL®, MySQL®, and MongoDB®. The procedures for adding these databases to your applications are the same in each case. MongoDB provides some extended management features; these are discussed in [Section 3.6.3.4, “Adding and Managing MongoDB Instances”](#).

3.6.3.1. Using Database Environment Variables

When you add a database cartridge, OpenShift automatically creates a database with the same name as the application and creates several environment variables. The following are examples of these variables:

```
OPENSIFT_DB_CTL_SCRIPT="/var/lib/libra//48afdc700f6240dabbb313f2a824f9ea//mysql-5.1//
wottle_mysql_ctl.sh"
OPENSIFT_DB_HOST="127.1.1.10"
OPENSIFT_DB_PASSWORD="ad3eR$4agk8z"
OPENSIFT_DB_PORT="3306"

OPENSIFT_DB_SOCKET="/var/lib/libra//48afdc700f6240dabbb313f2a824f9ea//mysql-5.1//socket/
mysql.sock"
```

```
OPENSIFT_DB_TYPE="mysql-5.1"
OPENSIFT_DB_URL="mysql://admin:F-7SCRSjAMZR@127.1.45.1:3306/"
OPENSIFT_DB_USERNAME="admin"
```

OpenShift does not currently support user changes to environment variables. This includes changing the default MySQL admin password (even outside of phpmyadmin). If you change your password, ensure that the change takes effect correctly.

This restriction applies to the default administrative user. You can add more users as required, and specify a password of your own choosing for these users.

3.6.3.2. Adding a MySQL Database

The following example illustrates how to add a MySQL database to the application "MyApp".

```
$ rhc app cartridge add -a MyApp -c mysql-5.1
Password:
Contacting https://openshift.redhat.com
Contacting https://openshift.redhat.com
API version: 1.1.1
Broker version: 1.1.1

RESULT:

Mysql 5.1 database added. Please make note of these credentials:

Root User: admin
Root Password: ad3eR$4agk8z
Database Name: MyApp

Connection URL: mysql://127.1.1.10:3306/
```

3.6.3.3. Adding a PostgreSQL Database

To use a PostgreSQL database with your application, use the same procedure as described above for adding a MySQL database, but use the PostgreSQL cartridge instead:

```
$ rhc app cartridge add -a MyApp -c postgresql-8.4
```

You can use SSH to connect to your application and verify the integrity of your database. Use the following command to connect to your application:

```
$ ssh UUID@appname-namespace.rhcloud.com
```

Use the password provided as output to the **rhc app cartridge add** command to connect to the database:

```
PGPASSWORD=password psql -h 127.0.251.129 -U admin -d appname
psql (8.4.9)
Type "help" for help.
appname=#
```

The last line in the output indicates a successful connection to the database. As a further check, you can run a simple query:

```
appname=# select 1 as a, 2 as b, 3 as c;
```

```
a | b | c
---+---+---
1 | 2 | 3
(1 row)
```

3.6.3.4. Adding and Managing MongoDB Instances

3.6.3.4.1. Overview

OpenShift supports **MongoDB**, a customizable back end database for web applications. OpenShift also supports **RockMongo**, a **MongoDB** administration tool. You can add **RockMongo** to your applications and use it to manage your **MongoDB** instances.

3.6.3.4.2. Using MongoDB with OpenShift Applications

Use the `rhc app cartridge add` command to add a **MongoDB** database instance to your application and to add **RockMongo** to control that database instance. The following procedure demonstrates how to create a new application and how to use **MongoDB** and **RockMongo** with that application.

Procedure 3.10. How to use MongoDB and RockMongo with OpenShift applications

1. Create a new application.

```
$ rhc app create -a myMongo -t php-5.3
```

2. Add the **MongoDB** cartridge to your application.

```
$ rhc app cartridge add -a myMongo -c mongodb-2.0
```

Take note of the credentials and other information that this command provides.

3. Add the **RockMongo** cartridge to your application.

```
$ rhc app cartridge add -a myMongo -c rockmongo-1.1
```

Take note of the credentials and other information that this command provides.

3.6.3.4.3. Managing MongoDB Instances with a Browser

After you have added a **MongoDB** database and the **RockMongo** cartridge to your application, you can navigate to and start exploring your database. Use the URL and RockMongo User and RockMongo Password credentials provided as output from the last step in the [Procedure 3.10, "How to use MongoDB and RockMongo with OpenShift applications"](#) procedure. In this example, the URL is `https://myMongo-myDomain.rhcloud.com/rockmongo/`.

The following illustration shows the **RockMongo** web interface:

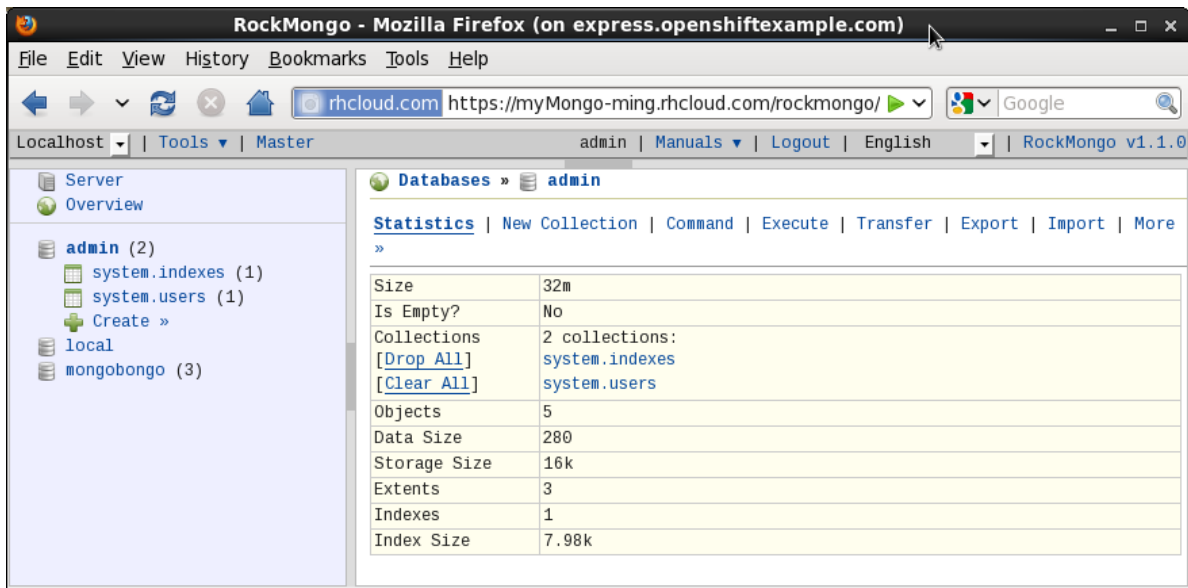


Figure 3.1. Example administration page of the RockMongo web interface

Refer to the **RockMongo** manuals for further information on how to use this interface to manage your database.

3.6.3.4.4. Managing MongoDB Instances in a Shell Environment

After you have added a **MongoDB** database and the **RockMongo** cartridge to your application as described in [Procedure 3.10, “How to use MongoDB and RockMongo with OpenShift applications”](#), you can manage your **MongoDB** instance in a shell environment.



Important

Shell access is quite powerful and it is possible to damage an application accidentally. Therefore it is recommended you only use shell access when necessary.

Use the **rhc domain show** command to display information about all of your available applications. Take note of the *UUID* and *Public URL* information that this command provides for your **MongoDB** instance.

The example below shows viewing information for the **myMongo** application:

Example 3.2. Viewing Application Information

```
$ rhc domain show
Password:
Application Info
=====
myMongo
  Framework: php-5.3
  Creation: 2011-12-07T01:07:33-05:00
  UUID: 0bd9d81bfe0a4def9de47b89fe1d3543
  Git URL: ssh://0bd9d81bfe0a4def9de47b89fe1d3543@myMongo-ming.rhcloud.com/~/.git/
  racer.git/
  Public URL: http://myMongo-ming.rhcloud.com/
```

Embedded:

```
rockmongo-1.1 - URL: https://myMongo-ming.rhcloud.com/rockmongo/
mongodb-2.0 - Connection URL: mongodb://127.1.2.129:27017/
```

You will also need the *Root username* and the *Root Password* for your **MongoDB** instance that was provided as output from Step 2 in [Procedure 3.10, “How to use MongoDB and RockMongo with OpenShift applications”](#).

When you have the necessary information, use the **ssh** command to open a shell environment for your **MongoDB** instance, substituting the *UUID* and *Public URL* parameters noted previously, as demonstrated below:

```
$ ssh 0bd9d81bfe0a4def9de47b89fe1d3543@myMongo-ming.rhcloud.com
```

When you have accessed your OpenShift application in the shell environment, you can type **help** at the shell prompt to get a list of the specialized shell commands.

3.6.3.4.4.1. Opening a Mongo Shell

From the shell prompt, use the following command to open an authenticated Mongo shell to run all Mongo shell commands, substituting the *Root Username* and *Root Password* obtained previously for your **MongoDB** instance:

```
[myMongo-ming.rhcloud.com~]\> mongo -u admin -p zVumUTBNKbXz admin
```

The example below demonstrates an open Mongo shell, and use of the **show users** Mongo shell command:

Example 3.3. Running Mongo Shell Commands

```
mongo -u admin -p zVMongoDB shell version: 2.0.2-rc1
connecting to: 127.0.250.129:27017/admin
umUTBNKbXz admin
show users
{
  "_id" : ObjectId("4ee55d39078e94193206e157"),
  "user" : "admin",
  "readOnly" : false,
  "pwd" : "aba43436961fbc6145261a12ed94b8f7"
}
```

Refer to the **MongoDB** website and documentation for more information on Mongo shell commands.

3.7. Creating Application Snapshots

You can use the **rhc app snapshot save** command to create a snapshot of your application and download it as a gzip archive. This also downloads any locally created log and other files, and maintains the directory structure that exists on the server. For example:

```
$ rhc app snapshot save -a ApplicationName --filepath Local/Path/To/Save/Tarball
```

Refer to [Section 5.1, “Backing up and Restoring Configuration and User Data”](#) for further information about this command.

3.8. Managing Applications

Use the **rhc app** command to manage your applications, view application status, and start, stop, restart, reload and destroy applications. You can also use it to list, add, or remove cartridges. Refer to the manual pages for more information and a full list of the available options. This command uses the following syntax:

```
$ rhc app <action> -a ApplicationName [Optional Arguments]
```

3.8.1. Using Application Management Commands

The following table describes the application management command options that you can use to manage an existing application in the cloud.

Table 3.4. Application management command argument options

Option	Details
start	Start an application in the cloud.
stop	Stop an application that is currently running in the cloud.
restart	Restart an application in the cloud.
reload	Reload an application in the cloud.
status	Display the current status of an application in the cloud.
destroy	Remove an application from the cloud.

3.8.2. Managing Applications in a Shell Environment

You can manage your OpenShift applications in a shell environment to perform specialized operations and general debugging. The shell access provides specialized tools for managing your applications.



Important

Shell access is quite powerful and it is possible to accidentally damage an application. Therefore it is recommended you only use shell access when necessary.

3.8.2.1. Opening a Shell Environment for an Application Node

Before you can access your application in a shell environment, ensure that you have completed the following:

- Create a domain as described in [Example 3.1, “Creating a new domain”](#). The examples in this section assume you have created the `ecs` domain.
- Create an application as described in [Section 3.4, “Creating Applications”](#). The examples in this section assume you have created the `"racer"` application.

After you have created your domain and your application, use the **rhc domain show** command to display information about all of the current applications:

Example 3.4. Viewing Application Information

```
$ rhc domain show
Password:
Application Info
=====
racer
  Framework: php-5.3
  Creation: 2011-12-07T01:07:33-05:00
  UUID: 0bd9d81bfe0a4def9de47b89fe1d3543
  Git URL: ssh://0bd9d81bfe0a4def9de47b89fe1d3543@racer-ecs.rhcloud.com/~/.git/
  racer.git/
  Public URL: http://racer-ecs.rhcloud.com/
```

Take note of the *UUID* and *Public URL* information for the application you wish to manage in the shell environment. Use the **ssh** command to open a shell environment for your application, substituting the *UUID* and *Public URL* parameters noted previously.

The example below demonstrates opening a shell environment to the "racer" application:

Example 3.5. Opening a Shell Environment for an Application Node

```
$ ssh 0bd9d81bfe0a4def9de47b89fe1d3543@racer-ecs.rhcloud.com
Warning: Permanently added 'racer-ecs.rhcloud.com,174.129.154.20' (RSA) to the list of
known hosts.

Welcome to OpenShift shell

This shell will assist you in managing openshift applications.

type "help" for more info.

[racer-ecs.rhcloud.com ~]\>
```



Note

The shell environment to access your OpenShift applications is protected and restricted with Security-Enhanced Linux (SELinux) policies.

When you have accessed your OpenShift application in the shell environment, you can type **help** at the shell prompt to get a list of the specialized shell commands. As mentioned previously, you can also use general Linux commands to perform routine operations in the shell environment.

3.8.3. Managing Applications in a Secure Shell Environment on Windows

Managing OpenShift applications in a Secure Shell (SSH) environment on Windows requires the PuTTY SSH and Telnet client, and PuTTYgen SSH key generator. This section provides instructions

on how you can use PuTTY to establish an SSH connection to your OpenShift applications. The instructions are provided in four easy steps, as outlined below:

- Step 1: Download PuTTY and PuTTYgen for Windows
- Step 2: Convert OpenSSH keys to PuTTY format
- Step 3: Locate application username and hostname
- Step 4: Establish SSH connection using PuTTY



Important

Secure Shell (SSH) access is quite powerful, and it is possible to accidentally damage an application. Therefore, it is recommended you only use SSH access when necessary.



Note

The screenshots shown in the instructions below were taken on Windows XP. Where necessary, the differences among other Windows operating systems have been noted.

It is assumed that the application you want to establish an SSH connection to has already been created. In the examples used in this section, an SSH connection will be established to the application named *testapp*.

Step 1: Download PuTTY and PuTTYgen

Go to <http://www.chiark.greenend.org.uk>¹. Download the executable files of the latest version of PuTTY and PuTTYgen to your desired directory. Alternatively, you can also download the PuTTY installer file that will install all required PuTTY packages on your computer.

Step 2: Convert OpenSSH Keys to PuTTY Format

By default, the pair of SSH keys generated when installing client tools on Windows are in OpenSSH format. The pair consists of a private key, named **id_rsa**, and a public key, named **id_rsa.pub**. To establish an SSH connection to your application using PuTTY, the private SSH key must be converted to PuTTY format. The pair of SSH keys generated during initial configuration of OpenShift client tools are stored in the following folders:

- **c:\Documents and Settings\user\.ssh** for Windows XP
- **c:\Users\user\.ssh** for Windows 7

Follow the instructions below to convert your private OpenSSH key to PuTTY format.

¹ <http://www.chiark.greenend.org.uk/~sgtatham/putty/>

1. Double-click **puttygen.exe** to launch PuTTYgen. If necessary, click **Run** in the *Security Warning* window. If you installed the PuTTY software with the installer file, click Start, point to All Programs, point to PuTTY, and then click **PuTTYgen**.
2. Click **Conversions**, then click **Import key** to select your private OpenSSH key, as shown in the figure below.

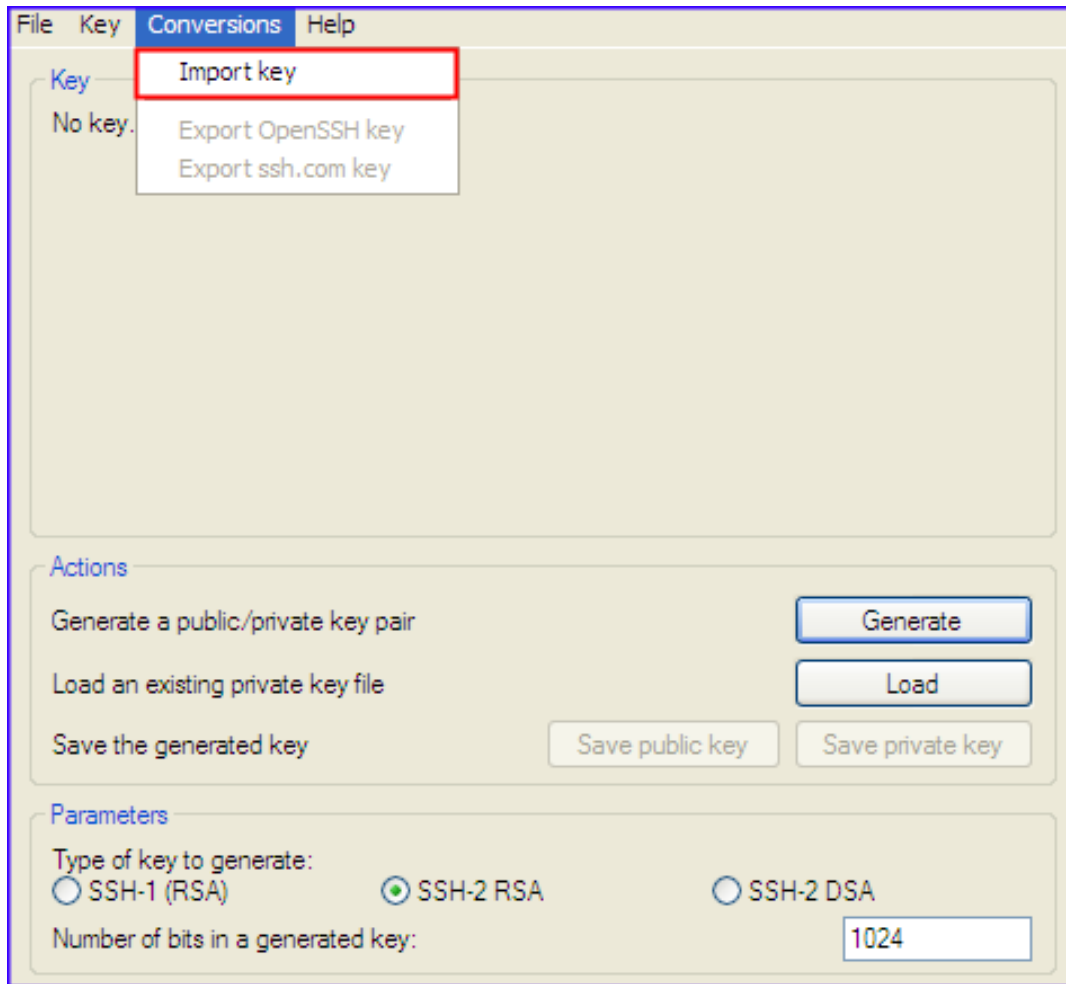


Figure 3.2. Import SSH Key

3. Navigate to the `\user\.ssh` folder, and select the `id_rsa` key file to import, as shown in the figure below.




Figure 3.3. Select SSH Key to Import

4. Your private SSH key should now be imported to PuTTYgen, as shown in the figure below. Click **Save private key**, and save the **id_rsa.ppk** file in the **\user\.ssh** folder where the original keys are stored.

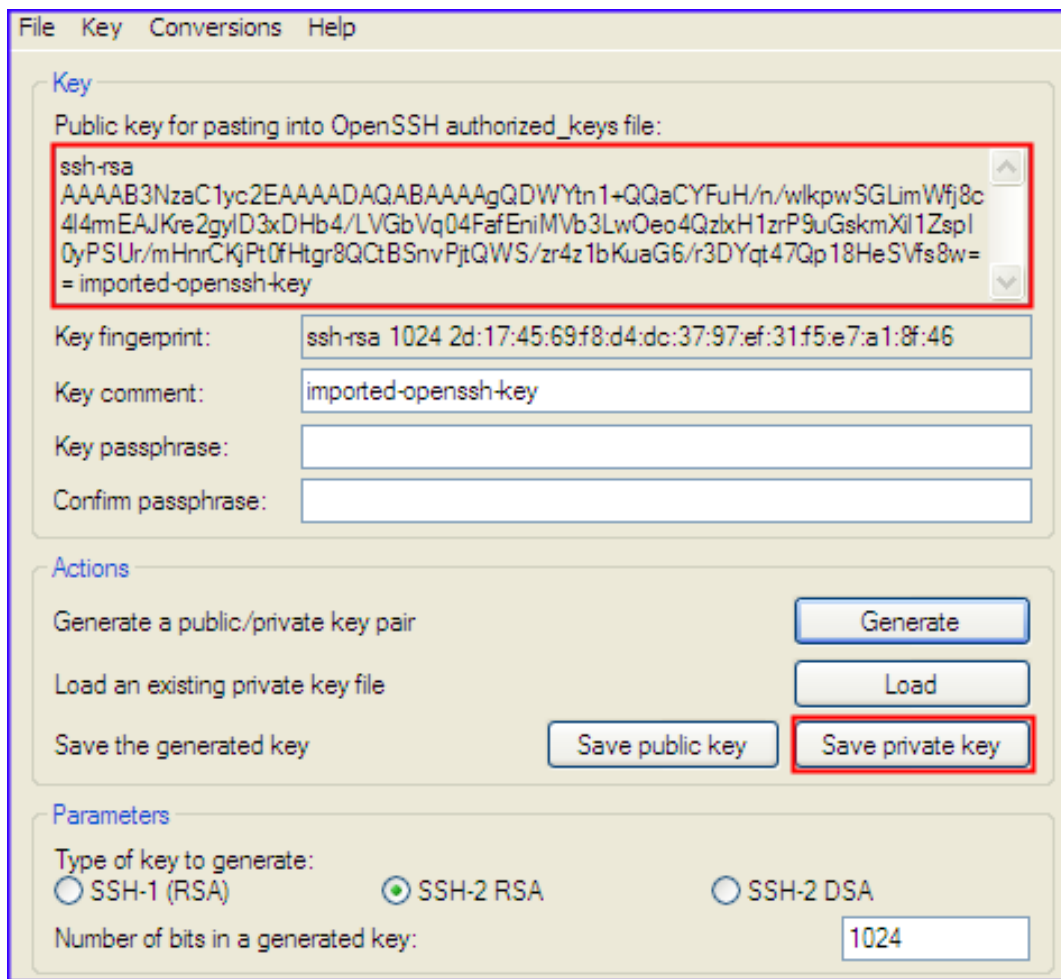


Figure 3.4. Save Key

5. Navigate to the `\user\.ssh` folder and verify that you now have three SSH key files. The `id_rsa.ppk` key will be used to establish an SSH connection to your OpenShift application. You can close PuTTYgen.

Step 3: Locate Application Username and Hostname

Each application created on OpenShift gets a unique UUID and gear name that is used to clone the remote Git repository to your local repository. You can use the `rhc domain info` command, or the OpenShift Management Console to get this information for the application to which you wish to establish an SSH connection. Both examples are shown below.

Use the `rhc domain info` command and note down the UUID and gear name of your application, as highlighted in the figure below. In this example the UUID and gear name are highlighted for the application named `testapp`.

```
testapp
  Framework: php-5.3
  Creation: 2012-06-18T01:02:39-04:00
  UUID: aa8d89ed311741e7b84d4edb82b11e0d
  Git URL: ssh://aa8d89ed311741e7b84d4edb82b11e0d@testapp-doctesting.rhcloud.com/~
  Public URL: http://testapp-doctesting.rhcloud.com/

Embedded:
  None
```

Figure 3.5 shows the output of the `testapp` command. The `Git URL` is highlighted in yellow, and the `UUID` is highlighted in green. A green box labeled "UUID / Username" points to the UUID, and a yellow box labeled "Gear name / Hostname" points to the gear name in the URL.

Figure 3.5. Application UUID and Gear Name

Alternatively, you can use the OpenShift Management Console to get the required information for your application. From the OpenShift Management Console, click on the **My Applications** tab, then click the name of the application you wish to access. The UUID and gear name can be found under the *GIT REPOSITORY* heading, as highlighted in the figure below.

Figure 3.6 shows the OpenShift Management Console interface for the **PHP 5.3** application. The `GIT REPOSITORY` heading is highlighted in green, and the `ssh://aa8d89ed311741e7b84d4edb82b11e0d@testapp-doctesting.rhcloud.com/~git/testapp.git/` URL is highlighted in yellow. A green box labeled "Username" points to the UUID, and a yellow box labeled "Hostname" points to the gear name in the URL.

Figure 3.6. Application UUID and Gear Name_WebUI

Take note of this information, because in the next step you will use the UUID as the username, and gear name as the hostname to configure PuTTY, and establish an SSH connection to your application. To avoid errors, it is recommended you cut and paste this information into PuTTY.

Step 4: Establish SSH Connection Using PuTTY

Now that you have the necessary information, you are ready to configure PuTTY to establish an SSH connection to your application. Follow the instructions below to configure PuTTY using the information from the previous step.

1. Double-click **putty.exe** to launch the PuTTY SSH and Telnet client. If necessary, click **Run** in the *Security Warning* window. If you installed the PuTTY software with the installer file, click **Start**, point to **All Programs**, point to **PuTTY**, and then click **PuTTY**.
2. In the left window pane under *Category*, click **Session**. Now copy and paste the gear name of your application in the *Host Name* text box, as highlighted in the figure below.

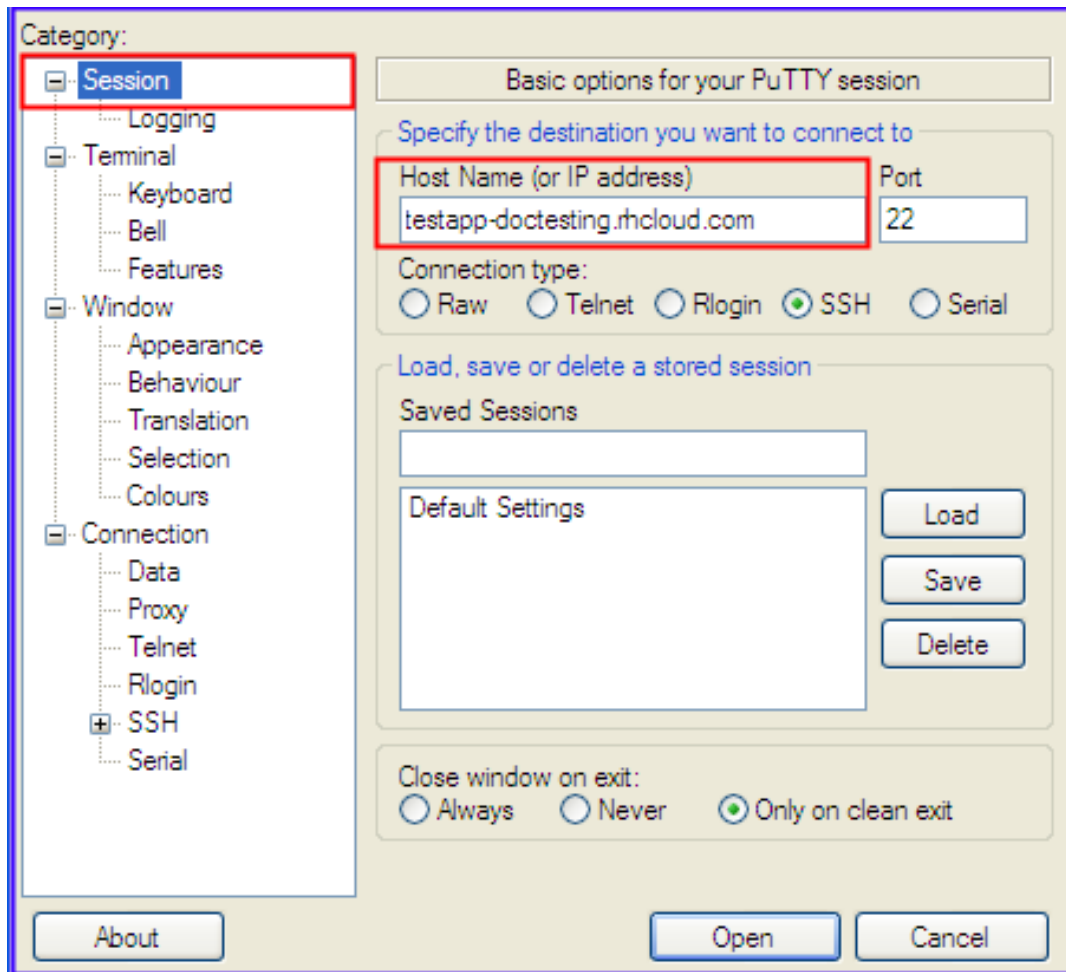


Figure 3.7. Enter Hostname in PuTTY

3. In the expanded list of options under **Connection**, click **Data**. Now copy and paste the UUID of your application in the *Auto-login username* text box, as highlighted in the figure below. Because the UUID is quite long, it may not be fully visible.

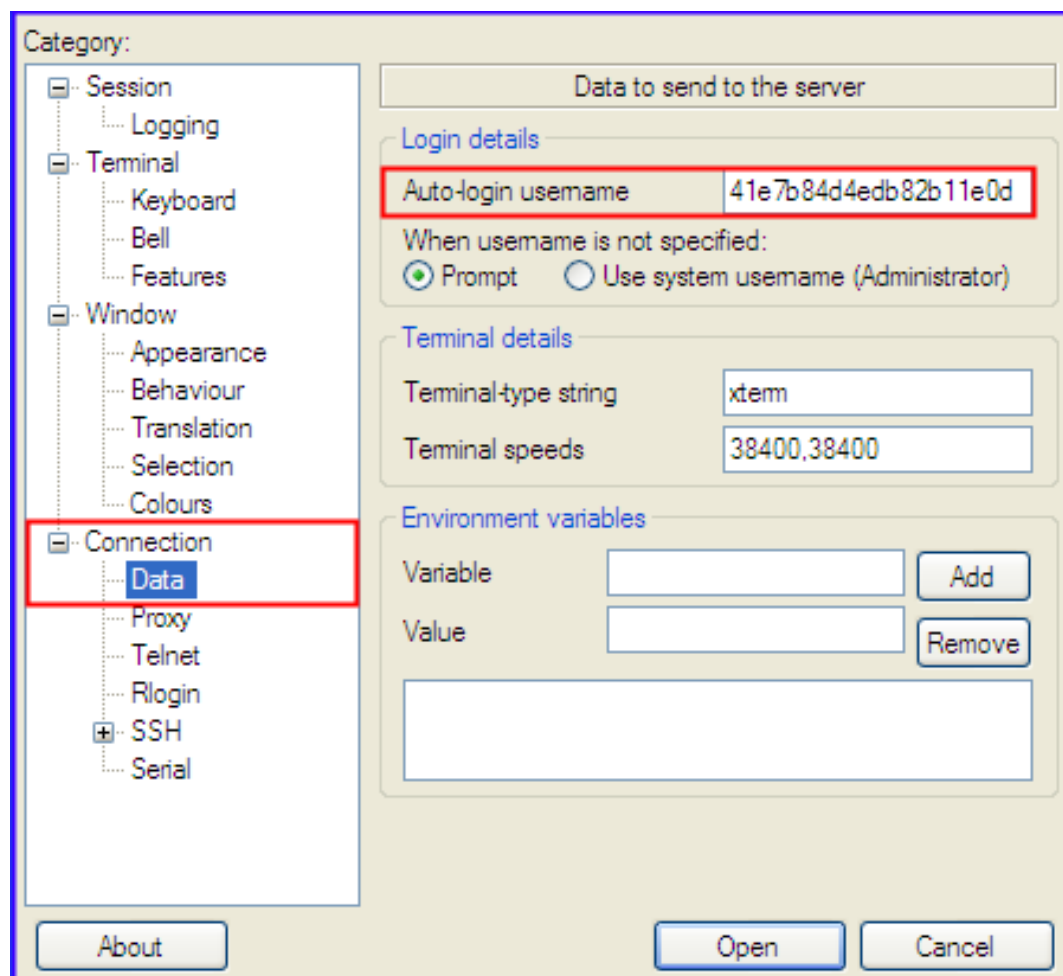


Figure 3.8. Enter UUID in PuTTY

4. In the expanded list of options under **Connection**|**SSH**, click **Auth**. Then click **Browse** to locate the **id_rsa.ppk** file you created earlier, as shown in the figure below.

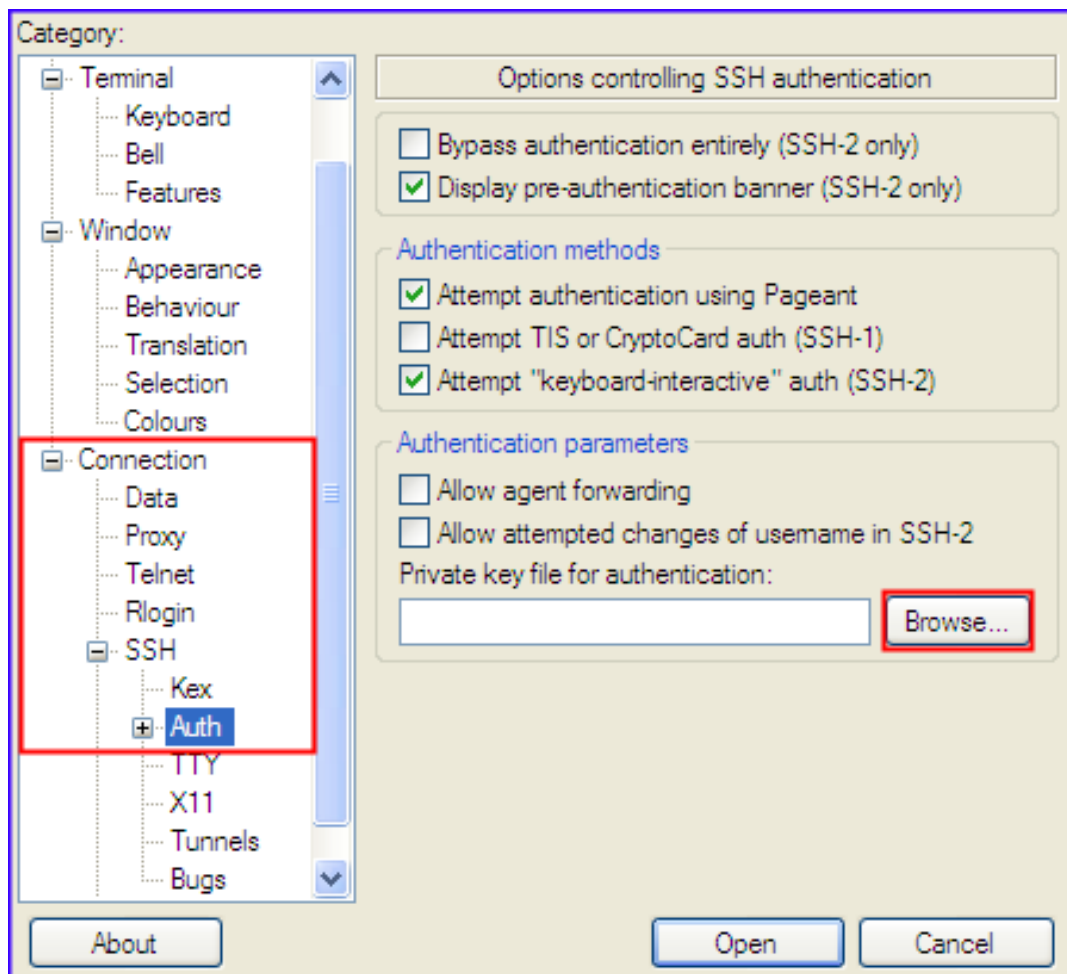


Figure 3.9. Find SSH Key

When you have browsed to the `id_rsa.ppk` file, click **Open**, as shown in the figure below.

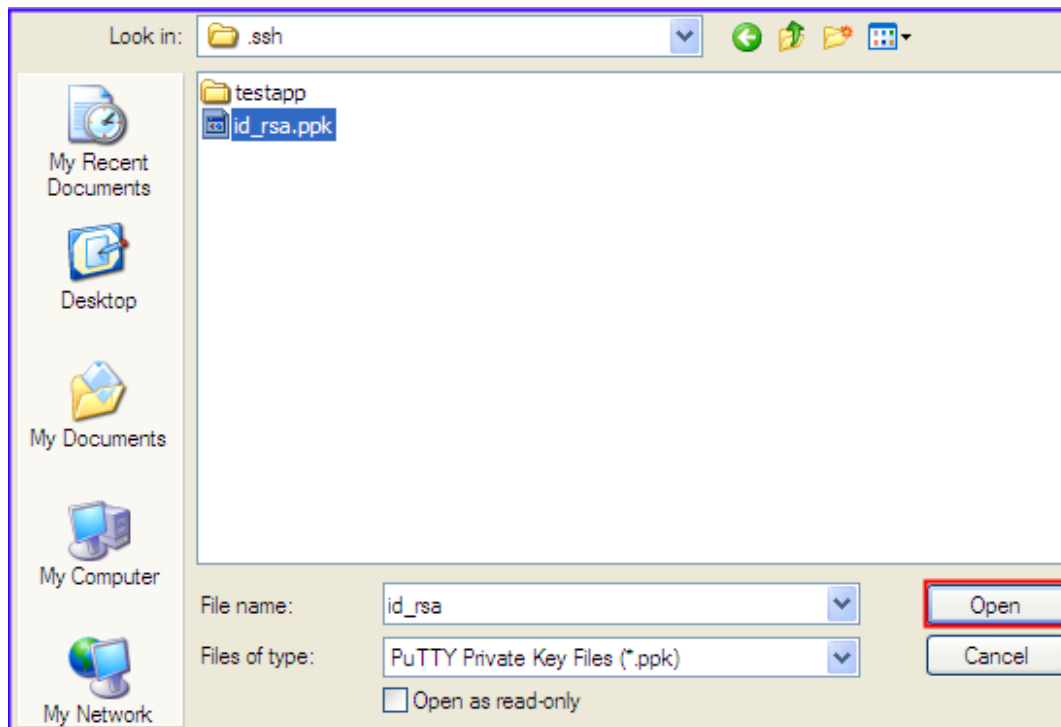


Figure 3.10. Select SSH Key

5. With your **id_rsa.ppk** key selected, click **Open**, as shown in the figure below. This will open an SSH connection to your application gear.

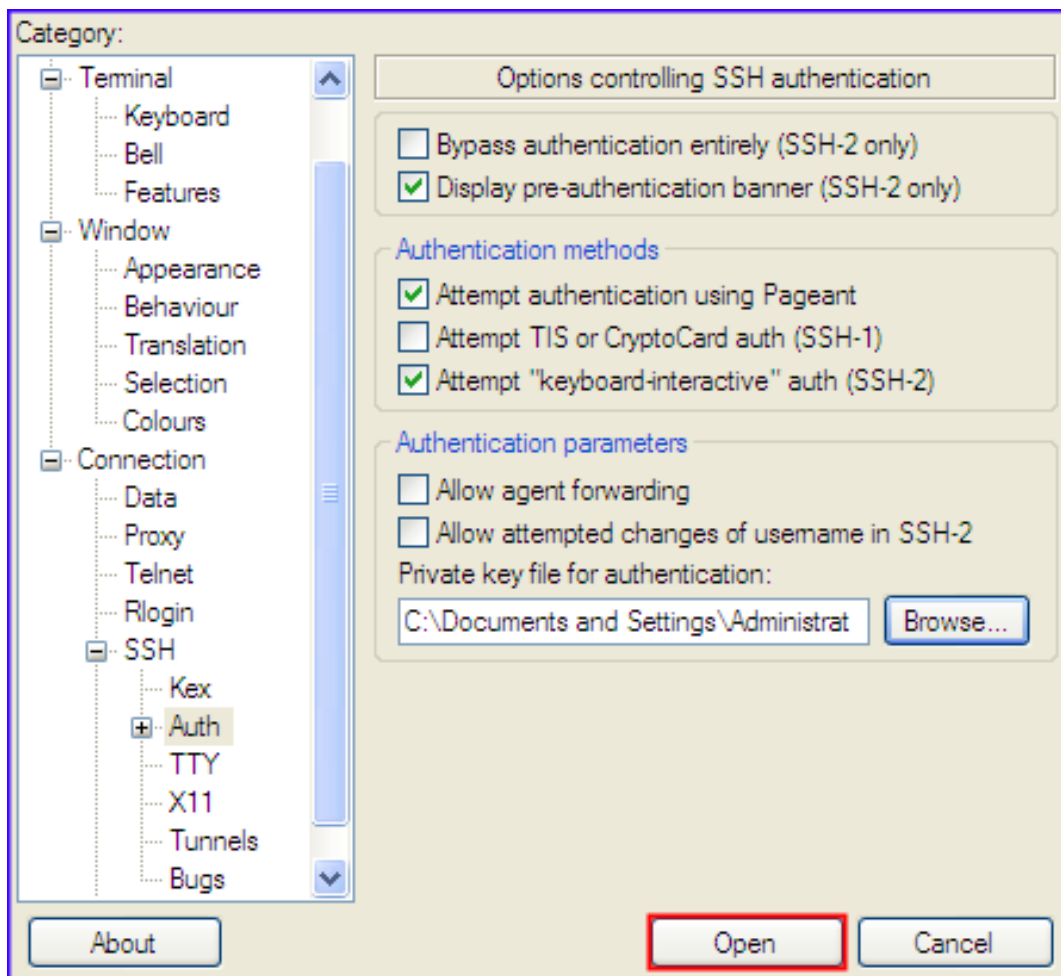


Figure 3.11. Connect to Application Gear

If a security warning is displayed, click **Yes** to accept it. An SSH terminal window will open, similar to the example shown below.

```
Using username "aa8d89ed311741e7b84d4edb82b11e0d".
Authenticating with public key "imported-openssh-key"

Welcome to OpenShift shell

This shell will assist you in managing OpenShift applications.

!!! IMPORTANT !!! IMPORTANT !!! IMPORTANT !!!
Shell access is quite powerful and it is possible for you to
accidentally damage your application. Proceed with care!
If worse comes to worst, destroy your application with 'rhc app destroy'
and recreate it
!!! IMPORTANT !!! IMPORTANT !!! IMPORTANT !!!

Type "help" for more info.

[testapp-doctesting.rhcloud.com ~]\>
```

One of the useful features of secure shell access is the ability to view application logs. You can find the application log files in the **App_Name/logs** directory. The example below shows how to view the log file for the application named *testapp*.

```
[testapp-doctesting.rhcloud.com ~]\> cd testapp/logs
[testapp-doctesting.rhcloud.com ~]\> ls
[testapp-doctesting.rhcloud.com ~]\> view error_log-20120618-000000-EST
```



Note

The Secure Shell environment to access your OpenShift applications is protected and restricted with Security-Enhanced Linux (SELinux) policies.

With an SSH connection established to your OpenShift application, you can type **help** at the shell prompt to get a list of specialized shell commands. You can also use general Linux commands to perform routine operations in the shell environment.

3.8.4. Using the Jenkins Embedded Build System

3.8.4.1. Introduction to Jenkins

Jenkins integrates with other OpenShift applications. To start building with Jenkins, you need to add the Jenkins Client to an existing application. From then on, running a **git push** command initiates a build process inside a Jenkins builder instead of inside your normal application compute space. For custom applications, or applications that have no upstream repositories, the build process can be initiated directly from the Jenkins web interface, without having to run the **git push** command.

Using the embedded Jenkins Client build system provides numerous benefits:

- Archived build information
- No application downtime during the build process
- Failed builds do not get deployed (leaving the previous working version in place)
- Jenkins builders have additional resources, for example memory and storage
- A large community of Jenkins plug-ins

3.8.4.2. The Build/Deploy Process in OpenShift

The OpenShift build/deploy process comprises four steps:

1. Pre-Receive - This occurs when you run a **git push** command, but before the push is fully committed.
2. Build - This builds your application, downloads required dependencies, executes the **.openshift/action_hooks/build** script and prepares everything for deployment.
3. Deploy - This performs any required tasks necessary to prepare the application for starting, including running the **.openshift/action_hooks/deploy** script. This step occurs immediately before the application is issued a **start** command.

4. Post-Deploy - This step allows for interaction with the running application, including running the `.openshift/action_hooks/post_deploy` script. This step occurs immediately after the application is started.

Building Without Jenkins

Building without Jenkins uses your application space as part of the build and test process. Because of this, the application is shut down so that its memory can be used for building.

Building with Jenkins

Building with Jenkins uses dedicated application space, which can be larger than the application runtime space. Because the build occurs in its own dedicated jail, the running application is not shut down or changed in any way until after the build succeeds. If the build fails, the currently active application will continue to run. However, a failure in the deployment process (deploy/start/post_deploy) may still leave the application partially deployed or inaccessible.

3.8.5. Managing Application Builds with Jenkins

3.8.5.1. Setting up Jenkins

This section describes how to set up Jenkins to automatically rebuild your application when you make changes to your local repository and then push those changes to the remote repository.



Note

The following procedures assume that you already have a valid user account for OpenShift.

3.8.5.1.1. Creating a Jenkins-enabled Application

New scaled or non-scaled applications can be enabled with Jenkins. You can use the **`rhc app create`** command to create a scaled or a non-scaled application with an associated Jenkins application, and to add the Jenkins client to your application.

Creating a non-scaled Jenkins-enabled application

Use the **`rhc app create`** command as shown below to create a non-scaled application with an associated Jenkins application, and to add the Jenkins client to your application.

Ensure that you take note of the login credentials that this command outputs to the screen. You will need these credentials to log in to the Jenkins home page.

```
$ rhc app create -a App01 -t php-5.3 --enable-jenkins
```

Creating a scaled Jenkins-enabled application

Use the **rhc app create** command with the `-s | --scaling` option to create a scaled application with an associated Jenkins application, and to add the Jenkins client to your application, as shown below:

```
$ rhc app create -a App01 -t php-5.3 --enable-jenkins -s
```

Ensure that you take note of the login credentials that this command outputs to the screen. You will need these credentials to log in to the Jenkins home page.

Confirming Your Jenkins-enabled Application

You can now run **rhc domain status** to confirm that your Jenkins-enabled application has been created. You can also navigate to the public URL returned by the **rhc app create** command to view your Jenkins home page.



Note

As part of the process of adding the Jenkins application to your domain, the Jenkins repository is downloaded to your local machine. This repository is not required for normal operations, and it is safe to delete this local copy.

3.8.5.2. Building Applications Without Jenkins

As described in [Building Without Jenkins](#), creating new builds of your application without using Jenkins requires that the currently running application be shut down in order to use its memory. This means that your application will be unavailable for the duration of the build and all associated tasks. These tasks are described below:

1. The user issues a **git push** command and the application is shut down.
2. The **git** application files are put in place.
3. **.openshift/action_hooks/build** is executed.
4. **.openshift/action_hooks/deploy** is executed.
5. **.openshift/action_hooks/post_deploy** is executed.
6. The application is restarted.

3.8.5.3. Building Applications with Jenkins

The Jenkins online build system monitors applications that have an embedded Jenkins Client, and will automatically rebuild and deploy those applications whenever changes to the git repository are pushed to the remote server. The existing application is not affected until a new, successful build has been created. Should the build fail, the existing application continues to run. No further interaction is required by the user.

The actual build and deployment process that Jenkins executes involves numerous steps, as described below.

1. The user issues a **git push** command, and Jenkins is notified that a new push is ready.

2. A dedicated Jenkins slave (a *builder*) is created. You can use the **rhc domain show** command to display this slave information. The application name is the same as the originating application, but with a "bldr" suffix.

**Important**

The first 28 characters of the application name must be unique or builders will be shared across applications. This can lead to build issues.

3. Jenkins runs the build.
4. Content from the originating application is downloaded to the builder application using **git** (for source code) and **rsync** (for existing libraries).
5. **ci_build.sh** is called from the Jenkins shell. This sets up the builder application for the Jenkins environment and performs some built-in bundling steps (PHP pear processing, Python virtual environment, etc).
6. **.openshift/action_hooks/build** is executed on the Jenkins builder.
7. Any additional desired steps are executed from the Jenkins shell (Maven build, Gem install, test cases, etc).
8. Jenkins stops the currently running application, and runs **rsync** to synchronize all new content over to the originating application.
9. **.openshift/action_hooks/deploy** is executed on the originating application.
10. Jenkins starts the originating application, and **.openshift/action_hooks/post_deploy** is executed on this application.
11. Jenkins archives all build artifacts for later reference.
12. After 15 minutes of idle time, the "build app" will be destroyed and will no longer appear in the output of the **rhc domain show** command. The build artifacts however, will still exist in Jenkins and can be viewed there.

You can monitor the build job using the Jenkins interface. The interface provides an extensive range of information about the current build, build history, artifacts, as well as plug-ins to graph, track, run tests and perform other operations. The following is an example build history of an application built using the embedded Jenkins Client.

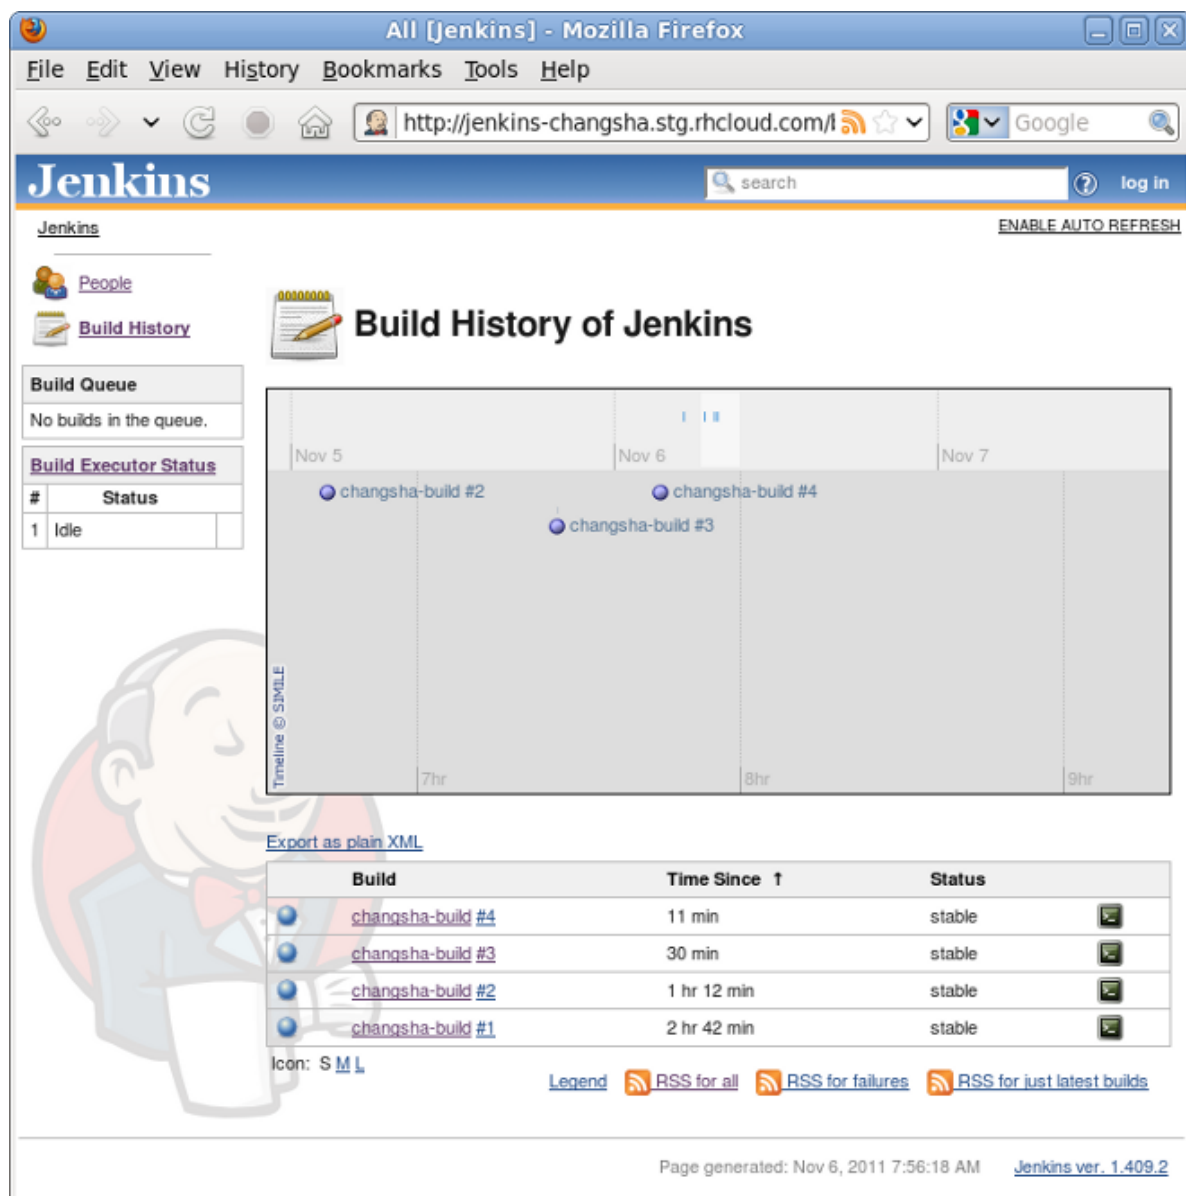



Figure 3.12. The Jenkins Build History Page

3.8.5.3.1. Building Custom Applications

You can also build custom applications, or applications that have no upstream repositories, directly from the Jenkins web interface instead of using the **git push** command.

To build your application from the Jenkins web interface, click on the  icon for the application you wish to build, located on the right side of the interface.

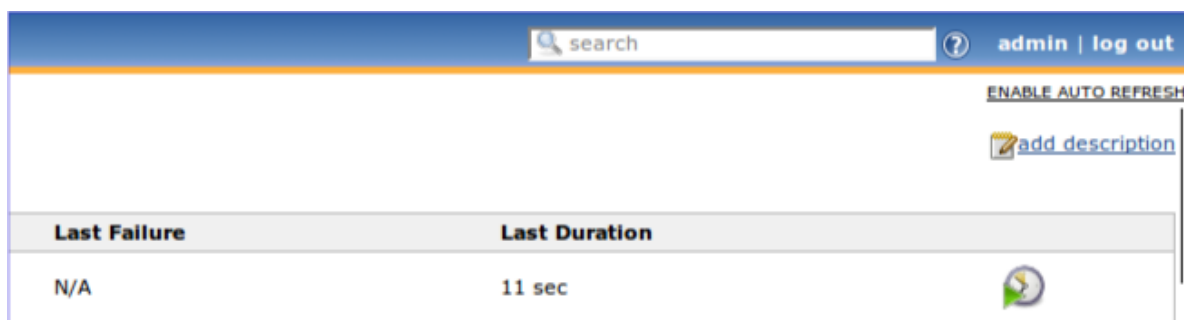


Figure 3.13. Build from Jenkins Interface

The status of the build process can be viewed in the web interface under the section labeled **Build Executor Status**.




Figure 3.14. Build Process Status

3.8.6. Using Environment Variables

Environment variables are named values that change the way running processes behave in computing environments. They can be used to store a variety of different values: application names, commonly accessed directory names, usernames, passwords, hostnames, and IP addresses.

OpenShift uses environment variables in the `/<appname>/openshift/config/standalone.xml` file that is part of `jbossas-7`. The following example is taken from this `standalone.xml` file, and contains three examples of included environment variables:

```
<connection-url>jdbc:mysql://${OPENSIFT_DB_HOST}:${OPENSIFT_DB_PORT}/
${OPENSIFT_APP_NAME}</connection-url>
```

For example, when you update your application using the **git push** command, OpenShift converts `OPENSIFT_DB_HOST` to the IP address of your database host.

You can also save environment variables on the server. This means that you do not have to pass sensitive information repeatedly to the command line.

Procedure 3.11. How to set environment variables on the server

1. Open the `<appname>/openshift/config/standalone.xml` file.
2. Specify the required values for any of your environment variables, and save and close the file.
3. Commit and push the changes to the server:

```
$ git commit -a -m "COMMIT MESSAGE"
$ git push
```

**Important**

Sensitive information stored in your environment variables is visible if you use the **rhc app snapshot** commands.

3.8.6.1. Common Environment Variables in OpenShift

The following are some common OpenShift environment variables. Note that this list is not exhaustive.

Table 3.5. Environment variables in OpenShift

Environment Variable Name	Use
System.getenv("OPENSIFT_APP_NAME")	Application name
System.getenv("OPENSIFT_GEAR_DIR")	Application directory
System.getenv("OPENSIFT_DATA_DIR")	For persistent storage (between pushes)
System.getenv("OPENSIFT_TMP_DIR")	Temporary storage (unmodified files are deleted after 10 days)

The following environment variables apply to the **rhc app** command when it is used to add a database:

Table 3.6. Environment variables applicable to database operations

Environment Variable Name	Use
System.getenv("OPENSIFT_DB_HOST")	Database host
System.getenv("OPENSIFT_DB_PORT")	Database port
System.getenv("OPENSIFT_DB_USERNAME")	Database username
System.getenv("OPENSIFT_DB_PASSWORD")	Database password

**Note**

To retrieve a complete list of environment variables, add an `export` statement to the `<appname>/openshift/action_hooks/build` script, and then run **git push**.

3.8.7. Using Node.js

This section provides the basic information required for running Node.js applications on OpenShift.

3.8.7.1. Repository Layout

When you create a new Node.js application on OpenShift, the application directory is populated with several files. The following table displays this layout:

Table 3.7. Node.js Repository Layout

File	Use
node_modules/	Node modules packaged with the application

File	Use
deplist.txt	List of npm modules to install with the application
npm_global_module_list	List of globally installed and available npm modules
server.js	OpenShift sample Node application
README	Information about using Node.js with OpenShift
.nodejs.env	Configuration file to set up environment and pass command line parameters to Node
.openshift/	Location for OpenShift specific files
.openshift/action_hooks/pre_build	Script that gets run every git push before the build
.openshift/action_hooks/build	Script that gets run every git push as part of the build process (on the CI system if available)
.openshift/action_hooks/deploy	Script that gets run every git push after build but before the application is restarted
.openshift/action_hooks/post_deploy	Script that gets run every git push after the application is restarted

You can create additional directories if needed, but do not delete the **node_modules** and **.openshift** directories.



Note

When you **git push**, everything in your remote repository directory is recreated. Place items that you want to persist between pushes, such as a database, in the **../data** directory.

3.8.7.2. Installing Node Modules

If you want to install Node modules from the [npm registry](http://npmjs.org)², you can specify the modules in the **deplist.txt** file. **Deplist.txt** contains a list of modules, one per line, with optional version information in the format: <module-name>@<version>. These modules install alongside your application ("locally") on the OpenShift environment when you **git push** your application.

Blank lines and lines starting with a hash (#) are ignored. To install sqlite3, for example, uncomment the line:

```
#sqlite3
```

Node modules not included in the [npm registry](http://npmjs.org)³ can be installed by packaging them in the **node_modules** directory.

The **npm_global_module_list** file contains a list of globally installed modules. You can override a globally available module by specifying it in **deplist.txt** or by packaging it in the **node_modules** directory. Node will give preference to the "locally" installed version of that module.

² <http://npmjs.org>

³ <http://npmjs.org>

3.8.7.3. Common Node.js Environment Variables in OpenShift

The following are some common OpenShift environment variables that you can use with your Node.js application. Note that this list is not exhaustive.

Table 3.8. Environment variables in OpenShift

Environment Variable Name	Use
process.env.OPENSIFT_GEAR_NAME	Application name
process.env.OPENSIFT_GEAR_DIR	Application directory
process.env.OPENSIFT_DATA_DIR	For persistent storage (between pushes)
process.env.OPENSIFT_TMP_DIR	Temporary storage (unmodified files are deleted after 10 days)

The following environment variables apply to the **rhc app** command when it is used to add a database:

Table 3.9. Environment Variables Applicable to Database Operations

Environment Variable Name	Use
process.env.OPENSIFT_DB_HOST	Database host
process.env.OPENSIFT_DB_PORT	Database port
process.env.OPENSIFT_DB_USERNAME	Database username
process.env.OPENSIFT_DB_PASSWORD	Database password

The following environment variables apply to the **rhc app** command when it is used to add a NoSQL database:

Table 3.10. Environment variables applicable to NoSQL database operations

Environment Variable Name	Use
process.env.OPENSIFT_NOSQL_DB_HOST	NoSQL database host
process.env.OPENSIFT_NOSQL_DB_PORT	NoSQL database port
process.env.OPENSIFT_NOSQL_DB_USERNAME	NoSQL database username
process.env.OPENSIFT_NOSQL_DB_PASSWORD	NoSQL database password



Note

To retrieve a complete list of environment variables, add an `export` statement to the `.openshift/action_hooks/build` script, and then run `git push`.

3.8.8. Scheduling Timed Jobs with Cron

You can use the OpenShift cron scheduler to set up timed jobs for your applications. This consists of adding the cron scheduler cartridge to your application, adding the required cron jobs to the appropriate directories, and then updating the remote git repository. You can also use the **rhc app** command to enable or disable your cron scripts.

The following example demonstrates how to create a new application and enable cron support.

Procedure 3.12. To create an application and enable cron support:

1. Create a new application:

```
$ rhc app create -a holy -t php-5.3
```

2. Add the cron scheduler cartridge to your new application:

```
$ rhc app cartridge add -a holy -c cron-1.4
```

3. Add the cron jobs to your application's **.openshift/cron/**
{minutely, hourly, weekly, daily, monthly}/ directories.

For example:

```
$ mkdir -p .openshift/cron/minutely
$ echo 'date >> $OPENSHIFT_REPO_DIR/php/date.txt' > .openshift/cron/minutely/date.sh
```

4. Commit your changes and push them to the remote repository.

```
$ git add .openshift/cron/
$ git commit -m "configuring cron jobs"
$ git push
```

This example appends a new line of date information to the **\$OPENSHIFT_REPO_DIR/php/date.txt** file every minute.

You can use **curl** to verify the operation of this script:

```
$ curl http://holy-roller.rhcloud.com/date.txt

Thu Feb  2 01:02:01 EST 2012
Thu Feb  2 01:03:01 EST 2012
Thu Feb  2 01:04:01 EST 2012
```

The scripts that you place in the **/cron** subdirectories are executed at the respective frequencies. Scripts in each subdirectory are executed sequentially, in alphabetical order. Scripts in the **/cron/hourly** directory are executed on the first minute of every hour. You can use the **rhc app cartridge** command to enable or disable your cron scripts. For example, to disable all of your scripts:

```
$ rhc app cartridge stop -a holy -c cron-1.4
```

To enable all of your scripts:

```
$ rhc app cartridge start -a holy -c cron-1.4
```



Note

The cron commands affect all cron jobs. You cannot disable or enable individual cron jobs.

3.8.9. Sending and Receiving E-mail

3.8.9.1. Overview

OpenShift provides support for connecting to externally hosted e-mail services (POP, IMAP, and SMTP). For developers, this means that you can establish a connection from your application to your own e-mail server, or to one of the popular public e-mail services, such as Gmail or YahooMail. If you are deploying popular blogging or wiki software, such as Drupal, Joomla, MediaWiki, or WordPress, on OpenShift, you can also take advantage of this feature by altering the e-mail settings of the software to point to the appropriate e-mail service.

3.8.9.2. E-mail Port Configuration

The specific ports that have been enabled to support this are as follows:

- SMTP/submission (25, 465, 587)
- IMAP (143, 220, 993)
- POP (109, 110, 995)

Communication occurs at a limited rate. Port 587 (submission) is restricted to a maximum rate of 256 kilobits per second. Ports 25 (SMTP) and 465 (SMTPS) are restricted to a maximum rate of 24 Kbps. Both will consume an extremely small share of the available bandwidth if congestion exists.



Important

Be aware that access to e-mail servers from cloud providers may be blocked by *Realtime Blackhole Lists (RBLs)*. This may affect your ability to connect to some e-mail servers. If you are unable to make a connection to one of these services, make sure that your e-mail provider allows authenticated connections from Amazon AWS EC2 hosts.

Application Maintenance, Monitoring and Troubleshooting

4.1. Monitoring Applications with the MongoDB Monitoring Service (MMS)

The MongoDB Monitoring Service (MMS) is a cloud-based monitoring service, designed by [10gen](https://www.10gen.com/)¹, to monitor the health of MongoDB deployments. MMS provides an agent that runs on MongoDB servers, and this agent reports information such as opcounters, memory information, number of connections, network I/O, database storage size, and more. All of this information is available in customizable charts that are displayed in your web browser.

OpenShift provides built-in support for MMS; you can add the MMS agent to your application using the same commands as for other cartridges. The only pre-requisite for using MMS is that you have an account with 10gen; you can sign up for a free MMS account at <https://mms.10gen.com/>.

4.1.1. Setting up MMS

The following procedure describes how to configure your application to take advantage of the services provided by MMS. This procedure assumes that you have successfully created an application and added the MongoDB cartridge.

Procedure 4.1. How to set up your application to use MMS

1. Download the agent.

When you log into MMS, you will see a "download the agent" link on the **Hosts** page. Click this link to download an agent preconfigured with your group credentials.

2. Create a directory named **mms** in your application's **.openshift** directory with the following command, replacing **app_directory** with the root directory for your application:

```
$ mkdir ~/app_directory/.openshift/mms
```

3. Add the **settings.py** file to the **mms** directory:

```
$ cp ~/mms-agent/settings.py ~/app_directory/.openshift/mms/
```



Important

The **settings.py** file contains the MMS agent settings that contain the API keys for connecting to your MMS account and updating the monitoring metrics. The MMS agent will not function without this file.

¹ <https://www.10gen.com/>

4. Use git to add and commit the **mms** directory to your local repository, then push it to your remote repository:

```
$ cd app_directory/  
$ git add .openshift/mms/  
$ git commit -m "mms settings" -a  
$ git push
```

5. Use the following command to add the agent to your application:

```
$ rhc app cartridge add -a app_name -c 10gen-mms-agent-0.1
```

After you have successfully completed this procedure, you can navigate to the <https://mms.10gen.com/> page, enter your host's details and start monitoring your application.

4.1.2. Monitoring Your Applications with MMS

4.1.2.1. Adding Hosts to MMS

After you have created an application and added the MMS agent, you can add the host to the **Hosts** page on <https://mms.10gen.com/>.² To add a new host, you need the hostname, port number, and login credentials that were provided when you added MongoDB to your application. If you no longer have this information, you can retrieve it directly from your application, as follows:

1. Use SSH to connect to your application:

```
$ ssh UUID@appname-namespace.rhcloud.com
```

where *UUID* is the UUID of your application. You can retrieve this using the **rhc domain show** command. Replace *appName* and *namespace* with your application name and namespace, respectively.

2. Run the following command to retrieve all the necessary connection and credential information for your application:

```
$ echo $OPENSIFT_NOSQL_DB_URL
```

Procedure 4.2. How to add hosts to MMS

1. Navigate to <https://mms.10gen.com/>
2. Click the **Hosts +** button at the top of the page to display the **New Host** dialog box.
3. Enter the required details and then click **Add**. This adds the host details to the agent, which immediately starts collecting and storing data.

² <https://mms.10gen.com/>

New Host [X]

MMS will automatically discover connected hosts. If you have a sharded cluster, you can seed MMS with a single mongos. If you have a replica set, you can add any member. If you have a master/slave cluster, you must add the slave(s).

Hostname

Port

Db Username

Db Password

(security credentials are optional)

Figure 4.1. Adding a new host to MMS

4.1.2.2. Using a Browser to Monitor Your Applications

After you have added the MMS agent to your application and added your host details to MMS, you can use the web interface to monitor your application's performance.

Procedure 4.3. How to monitor your applications with MMS

1. Navigate to <https://mms.10gen.com/> and click the **Hosts** tab.
2. Click the name of the host that you want to monitor to view the available data collection streams. You can customize this page to suit your own requirements. Refer to the 10gen documentation for full details.

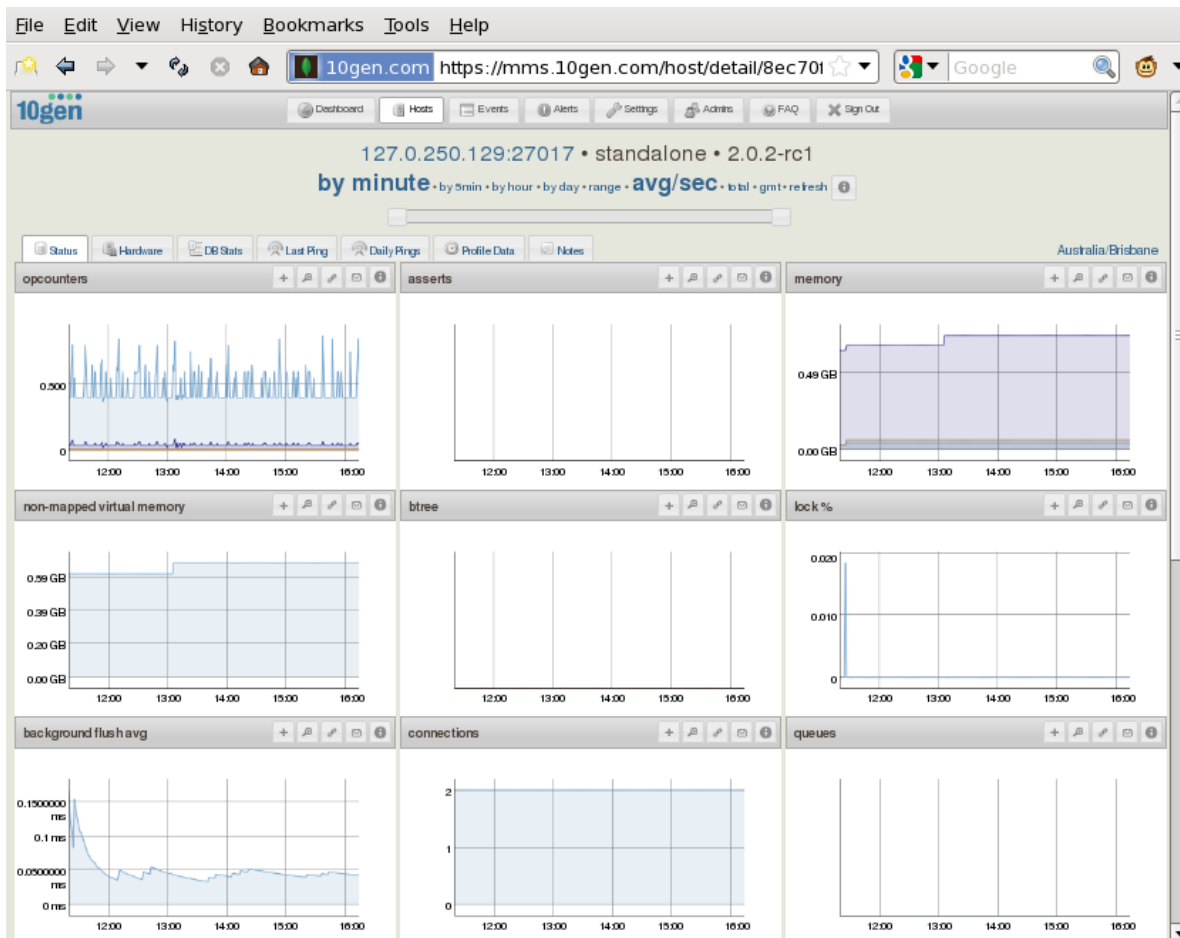


Figure 4.2. Example of data monitoring using MMS

4.2. Managing Your Application Disk Space

As you develop your applications and push changes to the git repository, you will slowly reduce the amount of disk space that is available to run your applications. This is because git stores all repository information, whether it is still required or not. Other aspects of developing and running your applications also result in wasted disk space, such as old log files and unused application libraries.

OpenShift provides tools to help you optimize your disk space to help achieve the best possible performance of your applications.

4.2.1. The Disk Space Cleanup Tool

OpenShift provides a *Disk Space Cleanup Tool* to help you manage your application disk space. This tool is part of the **rhc app** command suite, and performs the following functions:

- Runs the **git gc** command on your application's git repository on the server
- Clears the application's **/tmp** and log file directories. These are specified by the application's **OPENSHIFT_LOGS_DIR** and **OPENSHIFT_TMP_DIR** environment variables.

**Important**

OpenShift does not automatically back up or rotate log files. The Disk Space Cleanup Tool runs the **rm -rf** command to clear the contents of these directories. If you want to save their contents, you should use the **rhc app snapshot save** command first to create a snapshot of your system.

- Clears unused application libraries. This means that any library files previously installed by a **git push** command are removed.

The Disk Space Cleanup Tool uses the following syntax:

```
$ rhc app tidy -a <ApplicationName>
```

**Note**

The rhc client tools use the login specified in the `~/openshift/express.conf` file by default. If no default login is configured, or if you want to use a different login, include the `-l` parameter as a command line argument.

4.3. Troubleshooting JBoss Applications

OpenShift provides some basic tools to help you troubleshoot your JBoss applications should any problems arise. This includes the ability to inspect the log files and also to trigger a thread dump for JBoss AS applications.

4.3.1. Using Thread Dumps to Troubleshoot JBoss Applications

Thread dumps are a useful debugging tool. If an application appears to have stalled or is running out of resources, a thread dump can help to reveal the state of the server, to identify what might be causing any issues and ultimately to resolve the problem. You can use the **rhc app threaddump** command to trigger a thread dump for a JBoss AS application.

The following example demonstrates the use of this command:

**Note**

This example assumes that you have a valid account and have already created a domain and JBoss AS application called "myJBoss".

```
$ rhc app threaddump -a myJBoss
```

NOTE: The threaddump command is only applicable to applications of type jbossas-7. The thread dump file will be available via 'rhc app tail -f myJBoss/jbossas-7/stdout.log'

Password:

RESULT:

Success

You can now use the **rhc app tail** command to inspect the resultant thread dump:

```
$ rhc app tail -a myJBoss -f myJBoss/jbossas-7/stdout.log
Password:
Attempting to tail files: myJBoss/jbossas-7/stdout.log
Use ctrl + c to stop

Heap
PSYoungGen      total 5888K, used 3806K [0x00000000fe000000, 0x00000000fe5e0000,
  0x00000000100000000)
eden space 5760K, 65% used [0x00000000fe000000,0x00000000fe3af830,0x00000000fe5a0000)
from space 128K, 25% used [0x00000000fe5a0000,0x00000000fe5a8000,0x00000000fe5c0000)
to   space 128K, 0% used [0x00000000fe5c0000,0x00000000fe5c0000,0x00000000fe5e0000)
PSOldGen        total 64896K, used 9239K [0x00000000fa000000, 0x00000000fdf60000,
  0x00000000fe000000)
object space 64896K, 14% used [0x00000000fa000000,0x00000000fa905ce8,0x00000000fdf60000)
PSPermGen       total 33344K, used 33326K [0x00000000f4200000, 0x00000000f6290000,
  0x00000000fa000000)
object space 33344K, 99% used [0x00000000f4200000,0x00000000f628b930,0x00000000f6290000)

<output truncated>
```

4.3.2. Inspecting Server, Boot and Other Log Files

You can use the **rhc app tail** command to inspect other JBoss application log files. If you include your application name as the only argument, this command will tail the contents of all files in the `<appName>/logs/` directory. The following shows abbreviated output from this command:

```
$ rhc app tail -a myJBoss
Password:
Attempting to tail files: myJBoss/logs/*
Use ctrl + c to stop

==> myJBoss/logs/server.log <==
00:21:45,831 INFO  [org.apache.catalina.core.AprLifecycleListener] (MSC service thread 1-1)
  The Apache Tomcat Native library which allows optimal performance in production environments
  was not found on the java.library.path: /usr/lib/jvm/java-1.6.0-openjdk-1.6.0.0.x86_64/jre/
  lib/amd64/server:/usr/lib/jvm/java-1.6.0-openjdk-1.6.0.0.x86_64/jre/lib/amd64:/usr/lib/jvm/
  java-1.6.0-openjdk-1.6.0.0.x86_64/jre/..lib/amd64:/usr/java/packages/lib/amd64:/usr/lib64:/
  lib64:/lib:/usr/lib

00:21:46,569 INFO  [org.jboss.as.server.deployment] (MSC service thread 1-2) Starting
  deployment of "ROOT.war"
00:21:46,657 INFO  [org.jboss.as.jpa] (MSC service thread 1-2) added javax.persistence.api
  dependency to ROOT.war

==> myJBoss/logs/boot.log <==
sun.jnu.encoding = UTF-8
sun.management.compiler = HotSpot 64-Bit Tiered Compilers
user.country = US
user.dir = /tmp
00:21:45,337 INFO  [org.jboss.as.logging] Removing bootstrap log handlers

<output truncated>
```

The ability to obtain thread dumps and to inspect log files can greatly enhance your ability to troubleshoot your applications should any problems arise.

4.4. Performing Application Maintenance from Your Workstation

4.4.1. Port Forwarding

You can forward remote ports on your server to your workstation to make it easier to manage various services, such as MySQL. The **rhc-port-forward** command provides a wrapper for the **ssh** command, and determines which remote ports should be forwarded to your workstation.



Note

Ensure that your application is running before attempting to configure port forwarding.

Refer to the following example:

```
$ rhc-port-forward -a App01
Password:
Checking available ports...

Binding httpd -> 127.0.251.1:8080...

Use ctrl + c to stop
```

In this example, you could now open a browser and connect to your remote application as if it were running locally.

The current implementation of the **rhc-port-forward** command forwards all open ports on your running application to your local workstation. If you have multiple cartridges added to your application, you will see further messages indicating which remote services are being bound to local ports. Refer to the following example:

```
$ rhc-port-forward -a Pearl
Password:
Checking available ports...

Binding httpd -> 127.0.252.1:8080...
Binding mysqld -> 127.0.252.1:3306...

Use ctrl + c to stop
```

If you need to forward only specific ports, you need to use the **ssh** command with the **-L** option. Refer to the **ssh** manual page for details.

4.4.1.1. Port Forwarding on Mac OS X

At the time of this writing, out of the box Mac OS X only provides the following interfaces for loopback addresses:

- localhost

- 127.0.0.1

Therefore, you may experience error messages similar to those shown below when attempting to configure port forwarding using the IP address for your OpenShift application.

```
Checking available ports...
Binding httpd -> 127.10.51.129:8080...
Use ctrl + c to stop
bind: Can't assign requested address
channel_setup_fwd_listener: cannot listen to port: 8080
Could not request local forwarding.
```

The current workaround to enable port forwarding on Mac OS X is to manually configure an alias using the **ifconfig** command for each IP address used by your application, as shown below:

```
$ sudo ifconfig lo0 alias application_IP_address
```

For example, if the IP address used by your application is 127.10.51.129, run the command as shown below:

```
$ sudo ifconfig lo0 alias 127.10.51.129
```



Note

The **ifconfig** command on Mac OS X requires root/administrative privileges to execute.

You can now use the **rhc-port-forward** command to enable port forwarding.



Important

The IP address alias you configure for your OpenShift application is not persistent through system reboots. If you reboot your computer, you must repeat these steps to correctly enable port forwarding on Mac OS X.

Storage Management

OpenShift provides up to 512 megabytes of storage space for your application. You can also use your allocated storage space to store databases. The following table lists the currently supported database formats:

Table 5.1. Databases Supported by OpenShift

Database	Supported
MySQL	Yes
MySQLi	Yes
PDO MySQL	Yes
PDO SQLite	Yes
Amazon-RDS	Yes



Note

PHP Database Object (PDO) drivers are available only for MySQL and SQLite databases. If your application contains a database, place it in a new folder in your git repository. If you place the database in the `../data` folder, **git push** will not affect it.

5.1. Backing up and Restoring Configuration and User Data

Use the **rhc app snapshot save** command to create backups of your OpenShift applications, and also to restore those backups to the server. As a back-up tool, this command creates a gzipped tar file not only of your application but also of any locally-created log and other files, which is then downloaded to your local machine. The directory structure that exists on the server is maintained in the tar file.



Important

OpenShift does not maintain backups of your applications or user data on the OpenShift servers. The files created by this process are always downloaded to your local machine.

5.1.1. Creating Snapshots

The **rhc app snapshot save** command uses the following syntax to create a snapshot:

```
$ rhc app snapshot save -a Application_Name
```

The command will prompt for any missing information, such as your *rhlogin* and *password*. The default filename for the snapshot is ***\$Application_Name.tar.gz***. You can override this path and filename with the **--filepath** option.

The following example demonstrates the process of creating and downloading a snapshot:

```
$ rhc app snapshot save -a MasterApp
```

```
Password:
Pulling down a snapshot to MasterApp.tar.gz

Running extra dump: mysql_dump.sh
Mysql already running
72.7%
Stopping application...
Done
Creating and sending tar.gz
Running extra cleanup: mysql_cleanup.sh
Starting application...
Done
```

5.1.2. Restoring Snapshots

You can use the **rhc app snapshot restore** command to restore snapshots to the server. This form of the command restores the git repository, the application data directories and the log files found in the specified archive. When the restoration is complete, OpenShift runs the deployment script on the newly-restored repository, which completes the deployment process in the same way as if you had run **git push**.



Important

Even though you can save your application using a different filename, you cannot restore that application to a different name. That is, if your application name is "MyApp", you can save it as **OurApp.tar.gz**, but when you restore it you must use the original application name.

The **rhc app snapshot restore** command overwrites the remote git repository. Any changes that you might have made since you took the snapshot will be lost.

Appendix A. Revision History

Revision **Thu Aug 02 2012** **Brian Moss** docs@redhat.com
2.2.15-0

Add JBoss EAP 6.0 to application types that can hot deploy.
Add section on deploying JBoss applications on Java 6 or Java 7.
BZ842392: Clarify the application directory in the commands for setting up MMS.
BZ837493: Update procedure for generating new SSH keys.
BZ842517: Add instructions to change a domain name.
Update SSH authentication section.
Update list of preconfigured applications.
Update name relationship and platform container graphics.
BZ843332: Update section on adding an SSH key using Web UI.
Update cartridge information in Architecture Overview.

Revision **Tue Jul 10 2012** **Brian Moss** docs@redhat.com
2.2.14-0

Add Ruby 1.9 to list of available application types.
Update application types that can hot deploy.
Fixed BZ 837171: JBoss terminology.

Revision **Wed Jun 20 2012** **Bilhar Aulakh** docs@redhat.com
2.2.13-1

Add section on establishing an SSH connection using PuTTY on Windows.

Revision **Mon Jun 18 2012** **Brian Moss** docs@redhat.com
2.2.13-0

Add section on creating preconfigured applications.
Add JBoss EAP to list of available application types.
Add section on monitoring application resources.
Add section on hot deploying applications.
Move and edit scaling section.

Revision **Mon Jun 04 2012** **Brian Moss** docs@redhat.com
2.2.12-0

Fix typographical errors in section "Managing SSH Keys".
Update Preface with product and document information.
Document support for package.json.
Add MongoDB to available cartridges for scaled applications.
Reorganized cartridge content.
Updated Environmental Variable information.
Removed section on configuring local machine from WebUI chapter due to new configuration section in the Getting Started Guide.

Revision **Tue May 15 2012** **Brian Moss** docs@redhat.com
2.2.11-0

Appendix A. Revision History

Add Jenkins Server to list of application types and Jenkins Client 1.4 to list of cartridges for WebUI section.

Add content on using Node.js with OpenShift.

Edit webUI screen shots for updated web console.

BZ 821294: Missing namespace in Cron curl command.

BZ 821309: Minor QE fix.

Revision **Fri Apr 20 2012** **Bilhar Aulakh** docs@redhat.com
2.2.10-0

Spelling check.

Update section on timeout options for rhc app create command.

Update section on scaling, include support for jenkins-client-1.4 cartridges.

Add port forwarding on Mac OS X.

Revision **Fri Apr 13 2012** **Brian Moss** docs@redhat.com
2.2.9-0

BZ 811816: Update information on gears.

Revision **Tue Apr 10 2012** **Brian Moss** docs@redhat.com
2.2.8-3

BZ 809331: Minor QE fixes.

Add section on configuring local machine to WebUI chapter.

Revision **Wed Apr 04 2012** **David O'Brien** docs@redhat.com
2.2.8-2

Revise image tags to improve rendering.

Revision **Mon Apr 02 2012** **David O'Brien** docs@redhat.com
2.2.8-1

TA1698: Add information on gear sizes.

Revision **Fri Mar 30 2012** **Brian Moss** docs@redhat.com
2.2.8-0

Add screenshots to Web Interface Chapter.

Update images for Platform Overview section.

Update WebUI information for new Management Console.

Remove references to OpenShift Flex.

Remove references to "OpenShift Express" and replaced with "OpenShift".

Add section on scaled applications.

Add section on creating Jenkins-enabled scaled applications.

Revision **Tue Mar 20 2012** **Brian Moss** docs@redhat.com
2.2.7-0

Add sub-sections on architecture.

Renamed "raw" cartridge type to "Do it Yourself (DIY)".

Edits to WebUI Creating Applications section.
Moved Working with Namespaces section.

Revision **Thu Mar 01 2012** **Brian Moss** docs@redhat.com
2.2.6-1

Minor edits.
BZ 798626 - Update jbossas-7.0 references to jbossas-7
BZ 795238 - Remove section on Turkinator: no longer relevant.

Revision **Mon Feb 27 2012** **Brian Moss** docs@redhat.com
2.2.6-0

Add WebUI information.

Revision **Fri Feb 17 2012** **David O'Brien** docs@redhat.com
2.2.5-0

Add section on using e-mail ports.
Updated links to OpenShift web site and Bugzilla.
Document new rhc command syntax, improve examples.

Revision **Mon Feb 06 2012** **David O'Brien** docs@redhat.com
2.2.4-3

BZ 783103 - Improve section on creating applications.
BZ 783113 - Improved section on deploying applications.
BZ 787101 - Fix error in git command sequence.

Revision **Thu Feb 02 2012** **David O'Brien** docs@redhat.com
2.2.4-2

Add section on using cron for scheduled jobs.
Update supported python versions.

Revision **Wed Feb 01 2012** **Bilhar Aulakh** docs@redhat.com
2.2.4-1

Add section to manage SSH keys using web interface.
Edits to SSH key management section.

Revision **Fri Jan 20 2012** **David O'Brien** docs@redhat.com
2.2.4-0

BZ 783086 - Provide link to Bugzilla for reporting errors.

Revision **Fri Jan 13 2012** **Bilhar Aulakh** docs@redhat.com
2.2.3-5

BZ 781244 - Minor QE fixes.

Appendix A. Revision History

Revision **Fri Jan 13 2012** **David O'Brien** docs@redhat.com
2.2.3-4

Fix error in command example for destroying a domain.

Revision **Thu Jan 12 2012** **Bilhar Aulakh** docs@redhat.com
2.2.3-3

Add support for DSA SSH keys.

Add new CLI tools for multiple SSH key management.

Revision **Wed Jan 11 2012** **David O'Brien** docs@redhat.com
2.2.3-2

Add sub-section on using PostgreSQL.

BZ 773147 - Fix some typos and update rhc-user-info to rhc-domain-info.

Revise section on adding database cartridges.

Correct command on how to use alternative configuration file.

Revision **Fri Jan 06 2012** **David O'Brien** docs@redhat.com
2.2.3-1

Add section on how to trigger JBoss thread dumps.

Revision **Thu Dec 22 2011** **David O'Brien** docs@redhat.com
2.2.2-1

Add section on using 10gen MMS to monitor applications.

Revision **Mon Dec 19 2011** **David O'Brien** docs@redhat.com
2.2.2-0

Move to 3D version numbers.

Revision 2-1.3 **Thu Dec 15 2011** **Bilhar Aulakh** docs@redhat.com

Update usage of rhc-create-app for Jenkins.

Fix git command error in section on environment variables.

Add section on port forwarding.

Add section on managing MongoDB in a shell environment.

Revision 2-1.2 **Mon Dec 12 2011** **Bilhar Aulakh** docs@redhat.com

Add section on managing applications in a shell environment.

Revision 2-1.1 **Thu Dec 09 2011** **David O'Brien** docs@redhat.com

Add section on using environment variables.

Revision 2-1.0 **Tue Dec 06 2011** **David O'Brien** docs@redhat.com

Add section on using MongoDB and RockMongo.		
Revision 2-0.12	Fri Dec 02 2011	David O'Brien docs@redhat.com
Add section on using the Disk Space Cleanup Tool. Update section on Authenticating SSH Keys.		
Revision 2-0.11	Mon Nov 28 2011	Bilhar Aulakh docs@redhat.com
Minor grammar corrections.		
Revision 2-0.10	Fri Nov 25 2011	Bilhar Aulakh docs@redhat.com
Content and XML modifications.		
Revision 2-0.9	Fri Nov 25 2011	Bilhar Aulakh docs@redhat.com
Document the rhc-chk command.		
Revision 2-0.8	Thu Nov 24 2011	David O'Brien docs@redhat.com
Reworking section on rhc-create-domain.		
Revision 2-0.4	Mon Nov 22 2011	David O'Brien contentauthor@redhat.com
Update rhc-create-namespace command usage.		
Revision 2-0.3	Mon Nov 21 2011	David O'Brien contentauthor@redhat.com
Reword paragraph on how swap space works.		
Revision 2-0.0	Tue Nov 15 2011	David O'Brien docs@redhat.com
Moving to Edition 2.0.		
Revision 1-4.2	Tue Nov 15 2011	David O'Brien docs@redhat.com
Reset DPI on images to improve pdf display.		
Revision 1-4.1	Mon Nov 14 2011	David O'Brien docs@redhat.com
Minor update to restrictions on MySQL user passwords. Patch Express sign-up section to better match web site.		
Revision 1-4	Wed Nov 9 2011	David O'Brien docs@redhat.com
Added raw-0.1 to list of supported application types.		

Appendix A. Revision History

Update content for section "Using Arbitrary DNS Names" to improve example.

Revision 1-3 Wed Nov 9 2011

David O'Brien docs@redhat.com

Document availability of swap space for user applications.
Document support for custom domain names.
Updates to Jenkins sections.
Document Jenkins online building functionality.
Document restrictions on changing environment variables.
Add image to describe Express concepts, domain and username relationships.
Document adding applications using webUI.

Revision 1-2 Fri Oct 14 2011

David O'Brien docs@redhat.com

Document rhc-snapshot back-up and restore command.
Document "Maxretries = 7" when creating applications.
Document new embedded MySQL functionality.
Added jbossas-7.0 to list of supported application types.
Document -c option to specify config file for client tools.
Document blacklist check for domain name and app name.

Revision 1-1 Wed Aug 3 2011

David O'Brien docs@redhat.com

Minor corrections to deployment commands.
Further updates to repository layout.
Document existence of /tmp directory.
Revise section layout for JBoss AS application deployments.
Implement updates from tech review, improve examples.
Formatting updates, revised table layouts, spelling check.
Added section on implementing JBoss AS applications.

Revision 1-0 Fri Jun 10 2011

David O'Brien docs@redhat.com

Implement OpenShift branding.
Spell-check and formatting updates.
Update section on Turkinator.
Updates to Sign-up process.
Revise support and bug reporting options.

Revision 0-8 Tue May 03 2011

Misha Husnain Ali mhusnain@redhat.com

Final QE updates for Summit.

Revision 0-7 Thu Apr 29 2011

Misha Husnain Ali mhusnain@redhat.com

Further updates from QE and tech review.

Revision 0-6 Thu Apr 28 2011

Misha Husnain Ali mhusnain@redhat.com

Additional updates from initial QE and tech review.

Revision 0-5 Thu Apr 21 2011
Updates from initial QE and tech review.

David O'Brien contentauthor@redhat.com

Revision 0-4 Mon Apr 18 2011
Copy-edits and updates to procedures and terminology.

David O'Brien contentauthor@redhat.com

Revision 0-3 Sun Apr 10 2011
Content and XML edits

Misha Husnain Ali mhusnain@redhat.com

Revision 0-3 Wed Mar 30 2011
Updated sign-up process, application development and client tools installation

Muhammad Tahir mtahir@redhat.com

Revision 0-2 Thu Mar 24 2011
Minor XML updates and edits

Misha Husnain Ali mhusnain@redhat.com

Revision 0-1 Wed Mar 16 2011
Updated overview, user portal and installation notes

Muhammad Tahir mtahir@redhat.com

Revision 0-0 Thu Mar 3 2011
Initial creation of book by publican

Muhammad Tahir mtahir@redhat.com

Index

A

Applications
 alias, 25
 backing up
 command line, 37
 cloud
 git command, 28
 Web GUI, 15
 committing, 27
 creating, 24
 command, 25
 multiple, 25
 network issues, 24
 timeout, 24
 troubleshooting, 24
 Web GUI, 10, 10
 customized domain names
 alias, 25
 deleting
 command line, 38
 locally, 17
 Web GUI, 16
 deploying
 command line, 27
 destroying
 command line, 38
 directories, 1
 security, 1
 files, 12
 managing
 command line, 38
 preparing, 15
 reloading
 command line, 38
 removing
 command line, 38
 restarting
 command line, 38
 starting
 command line, 38
 stopping
 command line, 38
 testing, 27
 types
 no type, 27
 supported, 11
 viewing
 command line, 38

B

Build System

Jenkins, 50
 memory considerations, 51
 steps, 50

D

Do-It-Yourself, 11
 (see also Applications)

F

feedback
 contact information for this manual, viii

H

help
 getting help, viii

J

JBOSS AS, 11
 (see also Applications)
JBOSS EAP, 11
 (see also Applications)
Jenkins, 50
 (see also Build System)
 benefits, 50
 building
 space requirements, 51
Jenkins Server, 11
 (see also Applications)

N

Node.js, 11
 (see also Applications)

P

Perl, 11
 (see also Applications)
PHP, 11
 (see also Applications)
Python, 11
 (see also Applications)

R

Ruby Webserver Interface, 11, 11
 (see also Applications)

S

Security
 Environmental variables, 55
SSH, 9
 (see also Users)
Storage
 amount,

Index

types

- Amazon-RDS,
- MySQL,
- MySQLi,
- PDO MySQL,
- PDO SQLite,

U

Users

- managing

 - access, 9

- public keys, 9